

```
# QBank Scaffold v7:
# - Integrates full SH calibration into v5 backend
# - Adds RQ job runner + progress tracking
# - Admin UI progress screen
# - Unit tests for SH estimator (seeded synthetic pool)
# - Per-topic caps & blueprint-weighted  $\tau$ 
#
# Output: /mnt/data/qbank_scaffold_v7.zip
```

```
import os, zipfile, pathlib, textwrap, json
```

```
ROOT = "/mnt/data/qbank_scaffold_v7"
os.makedirs(ROOT, exist_ok=True)
```

```
def write(path, content):
    os.makedirs(os.path.dirname(path), exist_ok=True)
    with open(path, "w", encoding="utf-8") as f:
        f.write(textwrap.dedent(content))
```

```
# ----- README -----
```

```
write(f"{ROOT}/README.md", """
```

```
# QBank Scaffold (v7) — Integrated SH, Jobs, Progress UI, Topic Caps
```

What's new on top of v5b/v6:

- **Full Sympon–Hetter** iterative calibration integrated into the backend.
- **RQ job runner** (Redis) to run calibrations async, with progress updates.
- **Admin UI** page to launch calibrations and watch progress live.
- **Unit tests** for the SH estimator with a seeded synthetic pool.
- **Per-topic caps or blueprint-weighted τ** : calibrate using a scalar τ , a per-topic τ map, or topic weights.

```
## Quick start
```

```
1) Services (Postgres + Redis) up; DB schema from previous versions +
`item_exposure_control` present.
```

```
2) Backend API:
```

```
```bash
```

```
cd qbank-backend
```

```
pip install -r requirements.txt
```

```
uvicorn app.main:app --reload
```

```
Worker:
```

```
bash
```

```
Always show details
```

Copy  
cd qbank-backend  
python -m app.jobs.worker  
Admin UI:

bash  
Always show details

Copy  
cd admin-ui && npm install && npm run dev # http://localhost:4000  
Launch a calibration from the UI (Admin → Calibration) or via API:

bash  
Always show details

Copy  
POST /v1/admin/exposure/calibrate\_sh/start  
{  
 "exam\_code":"DEMO-EXAM", "tau":0.2, "n":400, "test\_len":25, "iters":5,  
 "alpha":0.6,  
 "topic\_tau": {"123":0.18, "124":0.22}, // optional  
 "topic\_weights": {"123":2.0, "124":1.0}, // optional  
 "dry\_run": false  
}

Cron / scheduled runs  
Kubernetes: see scripts/cron\_k8s/calibration-cronjob.yaml.

GitHub Actions (example only): .github/workflows/nightly\_calibration.yml.

Notes  
Progress is stored in RQ job meta: {current\_iter, total\_iters, avg\_exp, max\_over}  
and exposed via /status endpoint.

If both topic\_tau and topic\_weights are passed, topic\_tau takes precedence;  
otherwise weights are normalized to a  $\tau$  per topic.  
""")

----- Requirements -----  
write(f"{ROOT}/qbank-backend/requirements.txt", ""  
fastapi==0.115.0  
uvicorn==0.30.6  
pydantic==2.8.2  
python-dotenv==1.0.1

```
redis==5.0.8
rq==1.16.2
kafka-python==2.0.2
psycopg2-binary==2.9.9
SQLAlchemy==2.0.32
PyJWT==2.9.0
black==24.8.0
flake8==7.1.1
pytest==8.3.2
httpx==0.27.2
""")
```

----- Core config (add RQ settings) -----

```
write(f"{ROOT}/qbank-backend/app/core/config.py", """
import os
from dotenv import load_dotenv
load_dotenv()
```

```
DATABASE_URL = os.getenv("DATABASE_URL", "postgresql+psycopg2://
qbank:qbank@localhost:5432/qbank")
REDIS_URL = os.getenv("REDIS_URL", "redis://localhost:6379/0")
```

```
RQ
RQ_QUEUE = os.getenv("RQ_QUEUE", "calibration")
RQ_WORKER_CONCURRENCY = int(os.getenv("RQ_WORKER_CONCURRENCY",
"1"))
```

```
KAFKA_BOOTSTRAP = os.getenv("KAFKA_BOOTSTRAP", "localhost:9092")
KAFKA_TOPIC_EVENTS = os.getenv("KAFKA_TOPIC_EVENTS", "events.qbank")
TENANT_ID = os.getenv("APP_TENANT_ID",
"00000000-0000-0000-0000-000000000001")
APP_SECRET = os.getenv("APP_SECRET", "dev-secret-change-me")
MAX_DAILY_EXPOSURES = int(os.getenv("MAX_DAILY_EXPOSURES", "500"))
""")
```

```
write(f"{ROOT}/qbank-backend/app/core/cache.py", """
import redis
from datetime import datetime
from app.core.config import REDIS_URL, MAX_DAILY_EXPOSURES
```

```
redis_client = redis.Redis.from_url(REDIS_URL, decode_responses=True)
```

```
def exposure_key(question_id: int, version: int) -> str:
 day = datetime.utcnow().strftime("%Y%m%d")
```

```
return f"exp:{day}:{question_id}:{version}"
```

```
def can_serve(question_id: int, version: int) -> bool:
 key = exposure_key(question_id, version)
 count = int(redis_client.get(key) or 0)
 return count < MAX_DAILY_EXPOSURES
```

```
def bump_exposure(question_id: int, version: int) -> None:
 key = exposure_key(question_id, version)
 pipe = redis_client.pipeline()
 pipe.incr(key, 1)
 pipe.expire(key, 86400)
 pipe.execute()
 """
```

```
----- SH core (library) -----
```

```
write(f"{ROOT}/analytics/calibration/sh_core.py", """
import math, random, statistics
import psycpg2, psycpg2.extras
```

```
D = 1.7
```

```
def logistic(x: float) -> float: return 1.0 / (1.0 + math.exp(-x))
def prob_3pl(theta: float, a: float, b: float, c: float) -> float:
 return c + (1.0 - c) * logistic(D * a * (theta - b))
def fisher_info_3pl(theta: float, a: float, b: float, c: float) -> float:
 P = prob_3pl(theta, a, b, c); Q = 1.0 - P
 if P<=0 or Q<=0 or (1.0-c)<=0: return 0.0
 return (D2)*(a2)(Q/P)((P-c)/(1.0-c))**2
```

```
def sample_theta(dist: str = "normal0,1"):
 if dist.startswith("normal"): return random.gauss(0.0, 1.0)
 if dist.startswith("uniform"): return random.uniform(-1.0, 1.0)
 return 0.0
```

```
def load_pool(conn, exam_code: str):
 sql = '''
 SELECT qv.question_id, qv.version, qv.topic_id,
 COALESCE(ic.a, 1.0) AS a, COALESCE(ic.b, 0.0) AS b, COALESCE(ic.c, 0.2) AS c,
 COALESCE(iec.sh_p, 1.0) AS sh_p
 FROM question_publications qp
 JOIN question_versions qv ON qv.question_id=qp.question_id AND
 qv.version=qp.live_version
 LEFT JOIN item_calibration ic ON ic.question_id=qv.question_id AND
 ic.version=qv.version AND ic.model='3pl'
```

```
LEFT JOIN item_exposure_control iec ON iec.question_id=qv.question_id AND
iec.version=qv.version
WHERE qp.exam_code=%s AND qv.state='published'
'''
```

```
with conn.cursor(cursor_factory=psycopg2.extras.DictCursor) as cur:
cur.execute(sql, (exam_code,))
rows = cur.fetchall()
items = []
for r in rows:
items.append({
"qid": int(r["question_id"]), "ver": int(r["version"]),
"topic": r["topic_id"], "a": float(r["a"]), "b": float(r["b"]), "c": float(r["c"]),
"k": float(r["sh_p"]) if r["sh_p"] is not None else 1.0
})
return items
```

```
def simulate_once(pool, n: int, test_len: int, theta_dist: str, use_k: bool = True,
seed: int | None = None):
if seed is not None: random.seed(seed)
seen = { (it["qid"], it["ver"]): 0 for it in pool }
proposed = { (it["qid"], it["ver"]): 0 for it in pool }
```

python  
Always show details

Copy

```
for _ in range(n):
 theta = sample_theta(theta_dist)
 administered = set()
 for _pos in range(test_len):
 candidates = [it for it in pool if (it["qid"], it["ver"]) not in administered]
 if not candidates: break
 scored = [(fisher_info_3pl(theta, it["a"], it["b"], it["c"]), it) for it in
candidates]
 scored=sorted(key=lambda x: x[0], reverse=True)
 chosen = None
 for _, it in scored:
 key = (it["qid"], it["ver"])
 proposed[key] += 1
 if use_k:
 if random.random() <= max(0.0, min(1.0, it["k"])):
 chosen = it; break
 else:
 continue
```

```

 else:
 chosen = it; break
 if chosen is None: chosen = scored[0][1]
 key = (chosen["qid"], chosen["ver"])
 administered.add(key)
 for key in administered: seen[key] += 1
 return seen, proposed
def _compute_topic_tau(tau: float, topic_tau: dict | None, topic_weights: dict |
None) -> dict | None:
 if topic_tau: return {str(k): float(v) for k,v in topic_tau.items()}
 if topic_weights:
 # Normalize weights to sum to 1 and scale by tau
 s = sum(float(v) for v in topic_weights.values()) or 1.0
 return {str(k): tau * (float(v)/s) for k,v in topic_weights.items()}
 return None

```

```

def iterative_sh(pool, tau: float, n: int, test_len: int, iters: int, alpha: float,
theta_dist: str, floor: float, ceil: float, seed: int | None = None,
topic_tau: dict | None = None, topic_weights: dict | None = None):
 topic_tau_map = _compute_topic_tau(tau, topic_tau, topic_weights)
 k = { (it["qid"], it["ver"]): float(it["k"]) for it in pool }
 history = []
 for t in range(iters):
 for it in pool: it["k"] = k[(it["qid"], it["ver"])]
 seen, _ = simulate_once(pool, n=n, test_len=test_len, theta_dist=theta_dist,
use_k=True, seed=(None if seed is None else seed + t))
 r = { key: seen[key] / max(1, n) for key in seen }
 newk, diffs = {}, []
 for it in pool:
 key = (it["qid"], it["ver"])
 ri = rxget(key, 0.0); ki = k[key]
 # choose cap
 cap = tau
 if topic_tau_map is not None:
 cap = float(topic_tau_map.get(str(it["topic"]), tau))
 if ri <= 0.0:
 newk[key] = min(1.0, max(floor, ki * 1.1))
 else:
 ratio = cap / ri
 adj = ratio ** alpha
 cand = ki * adj
 newk[key] = min(1.0, max(floor, min(ceil, cand)))
 diffs.append(abs(ri - min(cap, ri)))
 k = newk

```

```

avg_exp = (sum(rxvalues())/len(r)) if r else 0.0
max_over = 0.0
if topic_tau_map is None:
max_over = max((ri - tau for ri in r.values()), default=0.0)
else:
compute per-topic overage
overages = []
for it in pool:
key = (it["qid"], it["ver"]); cap = float(topic_tau_map.get(str(it["topic"]), tau))
overages.append(r.get(key,0.0) - cap)
max_over = max(overages) if overages else 0.0
history.append({"iter": t+1, "avg_exp": avg_exp, "max_over": max_over})
return k, seen, history

```

```

def upsert_k(conn, kmap):
with conn.cursor() as cur:
cur.execute("SET search_path TO public")
for (qid, ver), kval in kmap.items():
cur.execute(
'''INSERT INTO item_exposure_control(question_id, version, sh_p)
VALUES (%s,%s,%s)
ON CONFLICT (question_id,version) DO UPDATE SET sh_p=EXCLUDED.sh_p,
updated_at=now()''',
(qid, ver, float(kval))
)
conn.commit()
'''')

```

```

----- CLI (wraps core) -----
write(f"{ROOT}/analytics/calibration/sh_iterative.py", """
import argparse, json, psycopg2
from .sh_core import load_pool, iterative_sh, upsert_k

```

```

def main():
ap = argparse.ArgumentParser()
ap.add_argument("--dsn", required=True)
ap.add_argument("--exam", "--exam_code", dest="exam_code", required=True)
ap.add_argument("--tau", type=float, default=0.2)
ap.add_argument("--n", type=int, default=500)
ap.add_argument("--len", "--test_len", dest="test_len", type=int, default=30)
ap.add_argument("--iters", type=int, default=6)
ap.add_argument("--alpha", type=float, default=0.6)
ap.add_argument("--theta", "--theta_dist", dest="theta_dist",
default="normal0,1")

```

```

ap.add_argument("--floor", type=float, default=0.02)
ap.add_argument("--ceil", type=float, default=1.0)
ap.add_argument("--seed", type=int, default=None)
ap.add_argument("--topic_tau", type=str, help="JSON {topic_id: tau_cap}",
default=None)
ap.add_argument("--topic_weights", type=str, help="JSON {topic_id: weight}",
default=None)
ap.add_argument("--dry_run", action="store_true")
args = ap.parse_args()

```

lua

Always show details

Copy

```

topic_tau = json.loads(args.topic_tau) if args.topic_tau else None
topic_weights = json.loads(args.topic_weights) if args.topic_weights else None

```

```

conn = psycopg2.connect(args.dsn)
pool = load_pool(conn, args.exam_code)
if not pool:
 print("No items found for", args.exam_code); return

```

```

kmap, seen, hist = iterative_sh(
 pool, tau=args.tau, n=args.n, test_len=args.test_len, iters=args.iters,
 alpha=args.alpha, theta_dist=args.theta_dist, floor=args.floor, ceil=args.ceil,
 seed=args.seed, topic_tau=topic_tau, topic_weights=topic_weights
)
print("history=", hist)
if not args.dry_run:
 upsert_k(conn, kmap)
 print("upserted=", len(kmap))
conn.close()
if name == "main":
 main()
"""

```

----- Jobs: RQ queue + worker + job function -----

```

write(f"{ROOT}/qbank-backend/app/jobs/queue.py", """
from rq import Queue
from redis import Redis
from app.core.config import REDIS_URL, RQ_QUEUE

```

```

redis = Redis.from_url(REDIS_URL)
queue = Queue(RQ_QUEUE, connection=redis)

```



```
""")
```

```
write(f"{ROOT}/qbank-backend/app/jobs/calibration_job.py", """)
import os, json, psycpg2
from rq import get_current_job
from analytics.calibration.sh_core import load_pool, iterative_sh, upsert_k

def calibrate_job(exam_code: str, dsn: str, tau: float, n: int, test_len: int, iters: int,
alpha: float, theta_dist: str, floor: float, ceil: float,
topic_tau: dict | None, topic_weights: dict | None, dry_run: bool):
 job = get_current_job()
 job.meta.update({"state": "running", "current_iter": 0, "total_iters": iters});
 job.save_meta()
```

```
python
Always show details
```

```
Copy
conn = psycpg2.connect(dsn)
pool = load_pool(conn, exam_code)
if not pool:
 job.meta.update({"state": "empty"}); job.save_meta(); conn.close(); return
{"updated": 0, "history": []}
```

```
iterative with callback-like progress (loop inside iterative_sh can't call back
easily, so call per-iter here)
kmap = { (it["qid"], it["ver"]): float(it["k"]) for it in pool }
history = []
for t in range(iters):
 # run one-iter by setting iters=1 and seeding changes each loop
 km, seen, hist = iterative_sh(pool, tau=tau, n=n, test_len=test_len, iters=1,
alpha=alpha,
 theta_dist=theta_dist, floor=floor, ceil=ceil, seed=None,
 topic_tau=topic_tau, topic_weights=topic_weights)
 # update pool ks
 for it in pool: it["k"] = km[(it["qid"], it["ver"])]
 kmap = km
 history.extend(hist)
 job.meta.update({"current_iter": t+1, "avg_exp": hist[-1]["avg_exp"],
"max_over": hist[-1]["max_over"]}); job.save_meta()
```

```
if not dry_run:
 upsert_k(conn, kmap)
conn.close()
```

```

job.meta.update({"state":"done"}); job.save_meta()
return {"updated": len(kmap), "history": history}
"""
)

```

```

write(f"{ROOT}/qbank-backend/app/jobs/worker.py", """
from rq import Worker
from app.jobs.queue import queue, redis
from app.core.config import RQ_QUEUE, RQ_WORKER_CONCURRENCY

```

```

if name == "main":
Note: RQ doesn't handle concurrency within a single worker process; scale by
processes
w = Worker([RQ_QUEUE], connection=redis)
w.work(with_scheduler=True)
"""
)

```

```

----- Admin API additions: start/status -----
write(f"{ROOT}/qbank-backend/app/api/admin.py", """
from fastapi import APIRouter, Depends, HTTPException
from pydantic import BaseModel, Field
from typing import List, Optional
from sqlalchemy.orm import Session
from sqlalchemy import text
import os, json
from rq.job import Job
from app.core.database import get_db
from app.core.auth import require_roles
from app.core.config import REDIS_URL
from app.jobs.queue import queue, redis
from app.jobs.calibration_job import calibrate_job
from app.models.orm import ItemExposureControl

```

```

router = APIRouter()

```

```

class ItemRow(BaseModel):
question_id: int
version: int
topic_id: Optional[int] = None
sh_p: float
recent_attempts: int | None = 0

```

```

@router.get("/exposure/items", response_model=List[ItemRow],
dependencies=[Depends(require_roles("admin"))])
def list_items(limit: int = 100, db: Session = Depends(get_db)):

```

```

rows = db.execute(text("""
SELECT qv.question_id, qv.version, qv.topic_id,
COALESCE(iec.sh_p, 1.0) as sh_p,
(SELECT count(*) FROM user_responses ur WHERE ur.question_id=qv.question_id
AND ur.version=qv.version AND ur.created_at>now()-interval '7 days') as
recent_attempts
FROM question_versions qv
LEFT JOIN item_exposure_control iec ON iec.question_id=qv.question_id AND
iec.version=qv.version
WHERE qv.state='published'
ORDER BY recent_attempts DESC NULLS LAST
LIMIT :lim
"""), {"lim": limit}).all()
return [ItemRow(question_id=r[0], version=r[1], topic_id=r[2], sh_p=float(r[3]),
recent_attempts=r[4] or 0) for r in rows]

```

```

class SetSh(BaseModel):
question_id: int
version: int
sh_p: float

```

```

@router.post("/exposure/set", dependencies=[Depends(require_roles("admin"))])
def set_sh(payload: SetSh, db: Session = Depends(get_db)):
if payload.sh_p < 0 or payload.sh_p > 1: raise HTTPException(400, "sh_p must be
in [0,1]")
row = db.get(ItemExposureControl, {"question_id": payload.question_id, "version":
payload.version})
if row: row.sh_p = payload.sh_p
else: db.add(ItemExposureControl(question_id=payload.question_id,
version=payload.version, sh_p=payload.sh_p))
db.commit()
return {"ok": True}

```

```

class StartCalib(BaseModel):
exam_code: str
tau: float = 0.2
n: int = 400
test_len: int = 25
iters: int = 5
alpha: float = 0.6
theta_dist: str = "normal0,1"
floor: float = 0.02
ceil: float = 1.0
topic_tau: Optional[dict] = None

```

```
topic_weights: Optional[dict] = None
dry_run: bool = False
```

```
@router.post("/exposure/calibrate_sh/start",
dependencies=[Depends(require_roles("admin"))])
def start_calibration(payload: StartCalib):
 dsn = os.getenv("DATABASE_URL", "postgresql+psycpg2://
qbank:qbank@localhost:5432/qbank")
 job = queue.enqueue(
 calibrate_job,
 payload.exam_code, dsn, payload.tau, payload.n, payload.test_len, payload.iters,
 payload.alpha, payload.theta_dist, payload.floor, payload.ceil,
 payload.topic_tau, payload.topic_weights, payload.dry_run,
 job_timeout=3600
)
 return {"job_id": job.get_id()}
```

```
class CalibStatus(BaseModel):
 state: str
 current_iter: int
 total_iters: int
 avg_exp: float | None = None
 max_over: float | None = None
 result: Optional[dict] = None
```

```
@router.get("/exposure/calibrate_sh/status", response_model=CalibStatus,
dependencies=[Depends(require_roles("admin"))])
def calib_status(job_id: str):
 job = Jobxfetch(job_id, connection=redis)
 meta = jobxmeta or {}
 state = metaxget("state") or job.get_status()
 return CalibStatus(
 state=state,
 current_iter=int(metaxget("current_iter") or 0),
 total_iters=int(metaxget("total_iters") or 0),
 avg_exp=float(metaxget("avg_exp") or 0.0) if meta.get("avg_exp") is not None
 else None,
 max_over=float(meta.get("max_over") or 0.0) if meta.get("max_over") is not None
 else None,
 result=job.result if state == "done" else None
)
 """
```

----- Main app (integrated) -----

```

write(f"{ROOT}/qbank-backend/app/main.py", """
from fastapi import FastAPI
from fastapi.middleware.cors import CORSMiddleware
from app.api.quizzes import router as quizzes_router
from app.api.author import router as author_router
from app.api.auth import router as auth_router
from app.api.admin import router as admin_router

app = FastAPI(title="QBank API v7", version="7.0.0")
app.add_middleware(CORSMiddleware, allow_origins=[""],
allow_credentials=True, allow_methods=[""], allow_headers=["*"])

app.include_router(auth_router, prefix="/v1/auth", tags=["auth"])
app.include_router(quizzes_router, prefix="/v1/quizzes", tags=["quizzes"])
app.include_router(author_router, prefix="/v1/author", tags=["authoring"])
app.include_router(admin_router, prefix="/v1/admin", tags=["admin"])

@app.get("/health")
def health(): return {"status": "ok"}
""")

----- Admin UI: calibration progress page -----
write(f"{ROOT}/admin-ui/pages/calibration.tsx", """
import { useEffect, useState } from 'react';
const API = process.env.NEXT_PUBLIC_API || 'http://localhost:8000';

type Status = { state:string; current_iter:number; total_iters:number;
avg_exp?:number; max_over?:number; result?:any };

export default function Calibration() {
const [token, setToken] = useState('');
const [jobId, setJobId] = useState<string>('');
const [status, setStatus] = useState<Status | null>(null);
const [form, setForm] = useState({ exam_code:'DEMO-EXAM', tau:0.2, n:400,
test_len:25, iters:5, alpha:0.6 });

const headers = { 'Content-Type':'application/json', 'Authorization': Bearer $
{token} };

const start = async () => {
const r = await fetch(`${API}/v1/admin/exposure/calibrate_sh/start,
{ method:'POST', headers, body: JSON.stringify(form) });
const data = await r.json(); setJobId(data.job_id);
};

```

```

useEffect(()=>{
const t = setInterval(async () => {
if (!jobId) return;
const r = await fetch(`${API}/v1/admin/exposure/calibrate_sh/status?job_id=${
jobId}, { headers });
if (r.ok) setStatus(await r.json());
}, 1500);
return () => clearInterval(t);
}, [jobId, token]);

const pct = status?.total_iters ? Math.round(100 * (status!.current_iter /
status!.total_iters)) : 0;

return (
<main style={{padding:24, maxWidth:800}}>
<h1>Calibration</h1>
<p>Paste an admin JWT</p>
<textarea value={token} onChange={(e)=>setToken(e.target.value)} rows={3}
style={{width:'100%'}} />
<section style={{marginTop:16}}>
<h3>Parameters</h3>
<div style={{display:'grid', gridTemplateColumns:'repeat(3, 1fr)', gap:12}}>
<label>Exam <input value={form.exam_code} onChange={(e)=>setForm({...form,
exam_code:e.target.value})} /></label>
<label> τ <input type="number" step="0.01" value={form.tau}
onChange={(e)=>setForm({...form, tau:parseFloat(e.target.value)})} /></label>
<label>n <input type="number" value={form.n} onChange={(e)=>setForm({...form,
n:parseInt(e.target.value)})} /></label>
<label>length <input type="number" value={form.test_len}
onChange={(e)=>setForm({...form, test_len:parseInt(e.target.value)})} /></label>
<label>iters <input type="number" value={form.iters}
onChange={(e)=>setForm({...form, iters:parseInt(e.target.value)})} /></label>
<label> α <input type="number" step="0.1" value={form.alpha}
onChange={(e)=>setForm({...form, alpha:parseFloat(e.target.value)})} /></label>
</div>
<button style={{marginTop:12}} onClick={start}>Start</button>
</section>
{status && (
<section style={{marginTop:24}}>
<h3>Status: {status.state}</h3>
<div style={{height:16, background:'#eee', borderRadius:8, overflow:'hidden'}}>
<div style={{width:${pct}%, height:'100%', background:'#4a90e2'}} />
</div>

```

```

<p style={{marginTop:8}}>iter {status.current_iter} / {status.total_iters} · avg_exp
{status.avg_exp?.toFixed(3)} · max_over {status.max_over?.toFixed(3)}</p>
{status.result && <pre style={{maxHeight:200, overflow:'auto',
background:'#fafafa', padding:12}}>{JSON.stringify(status.result, null, 2)}</pre>}
</section>
)}
</main>
);
}
""")

```

----- Unit tests for SH core -----

```

write(f"{ROOT}/analytics/tests/test_sh_core.py", """
from analytics.calibration.sh_core import iterative_sh

def test_iterative_sh_converges_under_tau():
Build a tiny synthetic pool with varied difficulty/discrimination
pool = [
{"qid":1,"ver":1,"topic":101,"a":1.2,"b":0.0,"c":0.2,"k":1.0},
{"qid":2,"ver":1,"topic":101,"a":0.8,"b":-0.5,"c":0.2,"k":1.0},
{"qid":3,"ver":1,"topic":102,"a":1.0,"b":0.5,"c":0.2,"k":1.0},
{"qid":4,"ver":1,"topic":102,"a":1.6,"b":0.0,"c":0.2,"k":1.0},
]
tau = 0.25
kmap, seen, hist = iterative_sh(pool, tau=tau, n=200, test_len=8, iters=4,
alpha=0.7, theta_dist="normal0,1", floor=0.02, ceil=1.0, seed=42)
final iteration summary must exist
assert hist[-1]["iter"] == 4
all k are within bounds
assert all(0.02 <= v <= 1.0 for v in kmap.values())
""")

```

----- Cron: K8s example & GitHub Actions example -----

```

write(f"{ROOT}/scripts/cron_k8s/calibration-cronjob.yaml", """
apiVersion: batch/v1
kind: CronJob
metadata:
name: qbank-calibration
spec:
schedule: "0 3 * * *"
jobTemplate:
spec:
template:
spec:

```

```

restartPolicy: Never
containers:
- name: calibrate
image: yourrepo/qbank-backend:latest
env:
- name: DATABASE_URL
valueFrom: { secretKeyRef: { name: qbank-secrets, key: DATABASE_URL } }
command: ["python","analytics/calibration/sh_iterative.py"]
args: ["--dsn","$(DATABASE_URL)","--exam","DEMO-EXAM","--tau","0.2","--n","1000","--len","30","--iters","6","--alpha","0.6"]
"""
)

```

```

write(f"{ROOT}/.github/workflows/nightly_calibration.yml", """
name: nightly-calibration
on:
schedule:
- cron: "0 4 * * *"
jobs:
run:
runs-on: ubuntu-latest
steps:
- uses: actions/checkout@v4
- uses: actions/setup-python@v5
with: { python-version: "3.11" }
- name: Install deps
run: pip install -r qbank-backend/requirements.txt
- name: Run calibration (requires network access to DB)
env:
DATABASE_URL: ${ secrets.QBANK_DATABASE_URL }
run: |
python analytics/calibration/sh_iterative.py --dsn "$DATABASE_URL" --exam
DEMO-EXAM --tau 0.2 --n 1000 --len 30 --iters 6 --alpha 0.6 --dry_run
"""
)

```

```

----- Zip -----
zip_path = "/mnt/data/qbank_scaffold_v7.zip"
with zipfile.ZipFile(zip_path, "w", zipfile.ZIP_DEFLATED) as z:
for p in pathlib.Path(ROOT).rglob("*"):
z.write(p, p.relative_to(ROOT))

```

zip\_path

Always show details



Copy

Result

'/mnt/data/qbank\_scaffold\_v7.zip'

Integrated and upgraded. Here's your bundle:

Download v7 (integrated SH + jobs + UI + tests)

What's in this cut:

Full SH inside your v5 backend

Refactor into a library: analytics/calibration/sh\_core.py (simulation + iterative updater + upsert).

CLI wrapper: analytics/calibration/sh\_iterative.py still available for ad-hoc/cron.

Async calibration via RQ

Queue/worker: qbank-backend/app/jobs/{queue.py, worker.py}.

Job fn updates progress in job.meta each iteration.

Admin API:

POST /v1/admin/exposure/calibrate\_sh/start → enqueues with params ( $\tau$ , n, len, iters,  $\alpha$ ... plus per-topic settings).

GET /v1/admin/exposure/calibrate\_sh/status?job\_id= → live progress (iter, avg exposure, max over-cap) + result on completion.

Progress UI for admins

admin-ui/pages/calibration.tsx: form to launch runs, progress bar with polling, result preview.

Per-topic caps & blueprint-weighted  $\tau$

Pass either topic\_tau (dict of topic\_id →  $\tau$ ) or topic\_weights (dict of weights). If both given, explicit caps win; otherwise weights are normalized to produce  $\tau$  per topic.

Unit tests

analytics/tests/test\_sh\_core.py: synthetic seed; verifies bounds & iteration flow.  
Extend easily with assertions on convergence if desired.

Scheduling examples

Kubernetes CronJob: scripts/cron\_k8s/calibration-cronjob.yaml.

GitHub Actions nightly (example/dry-run): .github/workflows/  
nightly\_calibration.yml.

Spin-up quick:

API + worker

```
bash
Copy
Edit
cd qbank-backend
pip install -r requirements.txt
uvicorn app.main:app --reload # terminal A
python -m app.jobs.worker # terminal B
Admin UI
```

```
arduino
Copy
Edit
cd admin-ui
npm install
npm run dev # http://localhost:4000/calibration
Launch a job (UI or cURL)
```

```
bash
Copy
Edit
POST /v1/admin/exposure/calibrate_sh/start
{
 "exam_code":"DEMO-EXAM","tau":0.2,"n":400,"test_len":25,"iters":5,"alpha":0.6,
 "topic_weights":{"101":2,"102":1}, "dry_run": false
}
```