

QBank v10 - Enterprise Question Bank System

Overview

QBank is a production-ready, enterprise-grade question bank platform designed for high-stakes educational assessment. Built with modern architecture principles, it provides advanced psychometric capabilities, adaptive testing, and comprehensive analytics.

Key Features

- **Advanced IRT Calibration:** Support for 1PL, 2PL, 3PL models with Bayesian estimation
- **Sympson-Hetter Exposure Control:** Iterative exposure control with probabilistic item selection
- **Adaptive Testing Engine:** Real-time ability estimation using EAP/MLE methods
- **Multi-tenant Architecture:** Complete tenant isolation with row-level security
- **Semantic Search:** Vector embeddings for intelligent question discovery
- **Real-time Analytics:** Streaming analytics with Kafka and ClickHouse
- **Enterprise Security:** OAuth2/SAML SSO, audit logging, RBAC
- **High Performance:** Async Python, Redis caching, horizontal scaling

Architecture



graph TB

subgraph "Frontend Layer"

A[Admin UI - Next.js]

B[Student Portal]

C[Author Portal]

end

subgraph "API Gateway"

D[FastAPI Backend]

E[GraphQL Gateway]

end

subgraph "Application Services"

F[Quiz Engine]

G[Adaptive Selector]

H[Calibration Service]

I[Analytics Service]

end

subgraph "Data Layer"

J[(PostgreSQL)]

K[(Redis Cache)]

L[(Elasticsearch)]

M[(ClickHouse)]

end

subgraph "Streaming"

N[Kafka]

O[Flink]

end

subgraph "ML Services"

P[IRT Engine]

Q[Recommender]

R[Embeddings]

end

A --> D

B --> D

C --> D

D --> F

D --> G

D --> H

```
D --> I
F --> J
F --> K
G --> P
H --> P
I --> M
F --> N
N --> O
O --> M
D --> L
G --> R
```

Prerequisites

- Docker & Docker Compose (v2.20+)
- Python 3.11+
- Node.js 18+
- PostgreSQL 15+ with pgvector extension
- Redis 7+
- Kafka 3.5+
- Elasticsearch 8+

Quick Start

1. Clone and Setup

```
bash

git clone https://github.com/yourorg/qbank.git
cd qbank
cp .env.example .env
# Edit .env with your configuration
```

2. Start Infrastructure

```
bash
```

```
# Start all services
```

```
docker-compose up -d
```

```
# Wait for services to be healthy
```

```
docker-compose ps
```

```
# Initialize database
```

```
docker-compose exec postgres psql -U qbank -d qbank -f /docker-entrypoint-initdb.d/01_core_schema.sql
```

```
docker-compose exec postgres psql -U qbank -d qbank -f /docker-entrypoint-initdb.d/02_content_ddl.sql
```

```
docker-compose exec postgres psql -U qbank -d qbank -f /docker-entrypoint-initdb.d/03_delivery_ddl.sql
```

```
docker-compose exec postgres psql -U qbank -d qbank -f /docker-entrypoint-initdb.d/04_analytics_ddl.sql
```

3. Run Migrations

```
bash
```

```
# Run Alembic migrations
```

```
docker-compose exec backend alembic upgrade head
```

```
# Seed initial data
```

```
docker-compose exec backend python scripts/seed_data.py
```

4. Access Services

- **API Documentation:** <http://localhost:8000/docs>
- **Admin UI:** <http://localhost:3000>
- **Kafka UI:** <http://localhost:8080>
- **Flower (Celery):** <http://localhost:5555>
- **Grafana:** <http://localhost:3001> (admin/admin)
- **Jaeger Tracing:** <http://localhost:16686>



API Documentation

Authentication

```
bash
```

Get access token

```
curl -X POST http://localhost:8000/v1/auth/login \  
-H "Content-Type: application/json" \  
-d '{"email": "admin@qbank.com", "password": "secure_password"}'
```

Use token in requests

```
curl -H "Authorization: Bearer <token>" \  
http://localhost:8000/v1/admin/dashboard/metrics
```

Core Endpoints

Questions & Content

python

Create question

POST /v1/author/questions

```
{  
  "external_ref": "MATH-001",  
  "topic_name": "Algebra",  
  "exam_code": "SAT-MATH",  
  "stem_md": "Solve for x:  $2x + 3 = 7$ ",  
  "lead_in": "What is the value of x?",  
  "rationale_md": "Subtract 3 from both sides...",  
  "difficulty_label": "medium",  
  "options": [  
    {"label": "A", "text_md": "1", "is_correct": false},  
    {"label": "B", "text_md": "2", "is_correct": true},  
    {"label": "C", "text_md": "3", "is_correct": false},  
    {"label": "D", "text_md": "4", "is_correct": false}  
  ]  
}
```

Search questions

GET /v1/questions/search?q=algebra&topic_id=1&difficulty=medium

Quiz Management

python

```
# Create adaptive quiz
POST /v1/quizzes
{
  "tenant_id": "00000000-0000-0000-0000-000000000001",
  "filters": {
    "topics": ["Algebra", "Geometry"],
    "difficulty": ["medium", "hard"],
    "num_questions": 30,
    "mode": "exam",
    "exam_code": "SAT-MATH"
  },
  "adaptive": true
}
```

```
# Get next question (adaptive)
GET /v1/quizzes/{quiz_id}/next
```

```
# Submit answer
POST /v1/quizzes/{quiz_id}/answers
{
  "question_id": 123,
  "selected": "B",
  "time_taken_ms": 45000,
  "confidence": 4
}
```

Calibration

```
python
```

```
# Start Simpson-Hetter calibration
POST /v1/admin/exposure/calibrate_sh/start
{
  "exam_code": "SAT-MATH",
  "tau": 0.2,
  "n": 1000,
  "test_len": 30,
  "iters": 10,
  "alpha": 0.8,
  "theta_dist": "normal(0,1)",
  "floor": 0.01,
  "ceil": 1.0
}

# Get calibration status
GET /v1/admin/exposure/calibrate_sh/runs/{run_id}
```

Testing

Unit Tests

```
bash

# Run all tests
docker-compose exec backend pytest

# Run with coverage
docker-compose exec backend pytest --cov=app --cov-report=html

# Run specific test
docker-compose exec backend pytest tests/test_adaptive.py::test_simpson_hetter
```

Load Testing

```
bash

# Install locust
pip install locust

# Run load test
locust -f tests/load/locustfile.py --host=http://localhost:8000
```

Analytics & Monitoring

Item Analysis

Access the analytics dashboard to view:

- P-values (difficulty indices)
- Discrimination indices
- Distractor analysis
- Response time distributions
- DIF (Differential Item Functioning) analysis

Real-time Metrics

Grafana dashboards provide:

- API response times
- Quiz completion rates
- User engagement metrics
- System health indicators
- IRT parameter distributions

Query Examples

```
sql
```



```
-- Item performance analysis
```

```
SELECT
```

```
  q.external_ref,  
  qv.difficulty_label,  
  is.p_value,  
  is.discrimination,  
  is.n_responses
```

```
FROM item_statistics is
```

```
JOIN question_versions qv ON is.question_id = qv.question_id  
  AND is.version = qv.version
```

```
JOIN questions q ON qv.question_id = q.id
```

```
WHERE is.p_value BETWEEN 0.3 AND 0.7  
  AND is.discrimination > 0.3
```

```
ORDER BY is.discrimination DESC;
```

```
-- Exposure analysis
```

```
SELECT
```

```
  question_id,  
  version,  
  sh_p,  
  exposure_count,  
  exposure_rate
```

```
FROM item_exposure_control
```

```
WHERE exposure_rate > 0.25
```

```
ORDER BY exposure_rate DESC;
```

Configuration

Environment Variables

Key configuration options in `.env`:

```
bash
```

Core Settings

ENVIRONMENT=production

DEBUG=false

SECRET_KEY=your-256-bit-secret-key

Database

DATABASE_URL=postgresql+asyncpg://user:pass@host:5432/dbname

DATABASE_POOL_SIZE=50

DATABASE_MAX_OVERFLOW=100

Redis

REDIS_URL=redis://host:6379/0

CACHE_TTL=3600

IRT Configuration

IRT_MODEL=3PL

IRT_MIN_RESPONSES=200

ADAPTIVE_ENABLED=true

Exposure Control

EXPOSURE_CONTROL_ENABLED=true

MAX_DAILY_EXPOSURES=500

DEFAULT_SH_P=1.0

ML Models

EMBEDDINGS_MODEL=sentence-transformers/all-MiniLM-L6-v2

DIFFICULTY_PREDICTION_ENABLED=true

Feature Flags

Control feature rollout via the database:

sql

```
-- Enable/disable features
```

```
UPDATE feature_flags
```

```
SET enabled = true,
```

```
value_json = '{"model": "3PL", "min_responses": 200}'
```

```
WHERE key = 'adaptive_testing';
```

```
-- Gradual rollout
```

```
UPDATE feature_flags
```

```
SET rollout_percentage = 50,
```

```
whitelist_users = '{"user1@example.com", "user2@example.com}"
```

```
WHERE key = 'new_calibration_algorithm';
```

Production Deployment

Kubernetes Deployment

```
bash
```

```
# Create namespace
```

```
kubectl create namespace qbank-prod
```

```
# Apply secrets
```

```
kubectl create secret generic qbank-secrets \
```

```
--from-env-file=.env.production \
```

```
-n qbank-prod
```

```
# Deploy application
```

```
kubectl apply -f k8s/ -n qbank-prod
```

```
# Check status
```

```
kubectl get pods -n qbank-prod
```

```
kubectl get services -n qbank-prod
```

Database Migrations

```
bash
```

Create migration

```
docker-compose exec backend alembic revision -m "Add new column"
```

Apply migrations (staging)

```
kubectl exec -it deployment/qbank-backend -n staging -- alembic upgrade head
```

Apply migrations (production)

```
kubectl exec -it deployment/qbank-backend -n production -- alembic upgrade head
```

Backup & Recovery

bash

Backup database

```
docker-compose exec postgres pg_dump -U qbank qbank | gzip > backup_$(date +%Y%m%d).sql.gz
```

Restore database

```
gunzip -c backup_20240115.sql.gz | docker-compose exec -T postgres psql -U qbank qbank
```

Backup Redis

```
docker-compose exec redis redis-cli BGSAVE
```

```
docker cp qbank-redis:/data/dump.rdb ./backups/redis_$(date +%Y%m%d).rdb
```



Performance Optimization

Database Optimization

sql

-- Analyze query performance

EXPLAIN ANALYZE

SELECT * FROM user_responses ur

JOIN question_versions qv ON ur.question_id = qv.question_id

WHERE ur.user_id = 'user123'

ORDER BY ur.created_at DESC

LIMIT 100;

-- Create partial indexes for common queries

CREATE INDEX CONCURRENTLY idx_ur_user_recent

ON user_responses(user_id, created_at DESC)

WHERE created_at > NOW() - INTERVAL '30 days';

-- Partition large tables

CREATE TABLE user_responses_2024_01 PARTITION OF user_responses

FOR VALUES FROM ('2024-01-01') TO ('2024-02-01');

Caching Strategy

python

Cache question pools

@cache.cache_result(prefix="pool", expire=3600)

async def get_question_pool(exam_code: str):

Expensive query

return await db.fetch_questions(exam_code)

Cache user abilities

@cache.cache_result(

prefix="ability",

expire=300,

key_builder=lambda user_id, topic_id: f"ability:{user_id}:{topic_id}"

)

async def get_user_ability(user_id: str, topic_id: int):

return await db.fetch_ability(user_id, topic_id)

Security

Authentication & Authorization

- JWT-based authentication with refresh tokens
- Role-based access control (RBAC)

- Row-level security (RLS) for multi-tenancy
- API rate limiting per user/IP

Data Protection

- Encryption at rest (PostgreSQL TDE)
- Encryption in transit (TLS 1.3)
- PII data segregation
- Audit logging for compliance

Security Headers

python

Configured in FastAPI middleware

```
security_headers = {  
    "X-Content-Type-Options": "nosniff",  
    "X-Frame-Options": "DENY",  
    "X-XSS-Protection": "1; mode=block",  
    "Strict-Transport-Security": "max-age=31536000; includeSubDomains",  
    "Content-Security-Policy": "default-src 'self'"  
}
```

Contributing

1. Fork the repository
2. Create a feature branch ((`git checkout -b feature/amazing-feature`))
3. Commit changes ((`git commit -m 'Add amazing feature'`))
4. Push to branch ((`git push origin feature/amazing-feature`))
5. Open a Pull Request

Code Style

bash

```
# Format code
black qbank-backend/
ruff check qbank-backend/

# Type checking
mypy qbank-backend/app

# Frontend linting
cd admin-ui && npm run lint
```



License

This project is licensed under the MIT License - see the [LICENSE](#) file for details.



Acknowledgments

- IRT algorithms based on research by Hambleton & Swaminathan
- Sympton-Hetter method from "A New Item Exposure Control Algorithm" (1985)
- FastAPI framework by Sebastián Ramírez
- UI components inspired by Tailwind UI



Support

- Documentation: <https://docs.qbank.io>
- Issues: <https://github.com/yourorg/qbank/issues>
- Discord: <https://discord.gg/qbank>
- Email: support@qbank.io

Built with ❤️ for the education community