

QBank Scaffold v3: Auth (JWT + RBAC), Alembic migration scripts,  
# exposure control + blueprint quotas in adaptive engine, extended ORM,  
# cohort analytics dbt models. Zips to /mnt/data/qbank\_scaffold\_v3.zip

```
import os, zipfile, pathlib, textwrap, json
```

```
ROOT = "/mnt/data/qbank_scaffold_v3"  
os.makedirs(ROOT, exist_ok=True)
```

```
def write(path, content):  
    os.makedirs(os.path.dirname(path), exist_ok=True)  
    with open(path, "w", encoding="utf-8") as f:  
        f.write(textwrap.dedent(content))
```

```
# ----- README -----  
write(f"{ROOT}/README.md", """)  
# QBank Backend Scaffold (v3)
```

New in v3:

- **Auth + RBAC**: JWT (HS256) with roles (`author`, `publisher`, `student`, `admin`).
- **Alembic migrations**: runnable SQL-first migrations.
- **Adaptive engine upgrades**: exposure control (per-item daily cap) + blueprint quota enforcement.
- **Extended ORM**: quiz sessions/items, responses, IRT calibration.
- **Cohort analytics (dbt)**: institution/tenant rollups.

> This is still a scaffold for rapid prototyping. Harden before production.

## Quick Start

1) Infra

```
```bash  
cd docker  
export APP_SECRET="change-me-please"  
docker compose up -d  
Initialize DB
```

bash

Always show details

Copy

```
psql postgresql://qbank:qbank@localhost:5432/qbank -f ../sql/content_ddl.sql  
psql postgresql://qbank:qbank@localhost:5432/qbank -f ../sql/delivery_ddl.sql  
psql postgresql://qbank:qbank@localhost:5432/qbank -f ../sql/analytics_ddl.sql
```

```
psql postgresql://qbank:qbank@localhost:5432/qbank -f ../sql/indexes.sql
psql postgresql://qbank:qbank@localhost:5432/qbank -f ../sql/auth_ddl.sql
(Or run Alembic: alembic -c ./alembic/alembic.ini upgrade head)
```

API

bash  
Always show details

Copy  
cd ../qbank-backend  
python -m venv .venv && source .venv/bin/activate  
pip install -r requirements.txt  
uvicorn app.main:app --reload  
# http://localhost:8000/docs  
Get a token (mock login)

bash  
Always show details

Copy  
curl -X POST http://localhost:8000/v1/auth/mock-login -H 'Content-Type: application/json' -d '{"user\_id":"user-123","roles":["author","publisher","student"]}'  
# use returned "access\_token" as "Authorization: Bearer <token>"

Notables  
Authoring endpoints now require roles:

POST /v1/author/questions → author role

POST /v1/author/publish/{question\_id} → publisher role

Adaptive selection respects:

Exposure control: Redis per-item cap (default 500/day; env MAX\_DAILY\_EXPOSURES)

Blueprint quotas: planned per-topic counts stored per quiz; selector favors topics with remaining quota

dbt: Adds cohort\_performance.sql by tenant and per-day.

""")

```
----- SQL DDL (adds auth) -----
write(f"{ROOT}/sql/content_ddl.sql", ""
CREATE EXTENSION IF NOT EXISTS ltree;
CREATE EXTENSION IF NOT EXISTS pgcrypto;
DO
B
E
G
I
N
P
E
R
F
O
R
M
1
F
R
O
M
p
g
e
x
t
e
n
s
i
o
n
W
H
E
R
E
e
x
t
n
a
m
e
=
```

*u  
e  
c  
t  
o  
r  
,  
;  
I  
F  
N  
O  
T  
F  
O  
U  
N  
D  
T  
H  
E  
N  
R  
A  
I  
S  
E  
N  
O  
T  
I  
C  
E  
,  
p  
g  
u  
e  
c  
t  
o  
r  
n  
o  
t  
i*

*n*  
*s*  
*t*  
*a*  
*l*  
*l*  
*e*  
*d*  
;  
*s*  
*k*  
*i*  
*p*  
,  
;  
*E*  
*N*  
*D*  
*I*  
*F*  
;  
*E*  
*N*  
*D*

BEGINPERFORM1FROMpg  
e

xtensionWHEREextname=  
,  
vector  
,  
;IFNOTFOUNDTHENRAISENOTICE  
,  
pgvectornotinstalled;skip  
,  
;ENDIF;END;

CREATE TABLE IF NOT EXISTS topics (  
id BIGSERIAL PRIMARY KEY,  
tenant\_id UUID NOT NULL DEFAULT  
'00000000-0000-0000-0000-000000000001',  
parent\_id BIGINT REFERENCES topics(id),  
name TEXT NOT NULL,  
blueprint\_code TEXT,  
path LTREE

);

```
CREATE TABLE IF NOT EXISTS questions (  
  id BIGSERIAL PRIMARY KEY,  
  tenant_id UUID NOT NULL DEFAULT  
  '00000000-0000-0000-0000-000000000001',  
  external_ref TEXT UNIQUE,  
  created_by UUID NOT NULL,  
  created_at TIMESTAMPTZ NOT NULL DEFAULT now(),  
  is_deleted BOOLEAN NOT NULL DEFAULT FALSE  
);
```

```
CREATE TABLE IF NOT EXISTS question_versions (  
  id BIGSERIAL PRIMARY KEY,  
  question_id BIGINT NOT NULL REFERENCES questions(id),  
  version INT NOT NULL,  
  state TEXT NOT NULL CHECK (state IN  
  ('draft','review','approved','published','retired')),  
  stem_md TEXT NOT NULL,  
  lead_in TEXT NOT NULL,  
  rationale_md TEXT NOT NULL,  
  difficulty_label TEXT,  
  bloom_level SMALLINT,  
  topic_id BIGINT REFERENCES topics(id),  
  tags JSONB NOT NULL DEFAULT '{}'::jsonb,  
  assets JSONB NOT NULL DEFAULT '[]'::jsonb,  
  references JSONB NOT NULL DEFAULT '[]'::jsonb,  
  created_by UUID NOT NULL,  
  created_at TIMESTAMPTZ NOT NULL DEFAULT now(),  
  UNIQUE(question_id, version)  
);
```

```
CREATE TABLE IF NOT EXISTS question_options (  
  id BIGSERIAL PRIMARY KEY,  
  question_version_id BIGINT NOT NULL REFERENCES question_versions(id) ON  
  DELETE CASCADE,  
  option_label CHAR(1) NOT NULL,  
  option_text_md TEXT NOT NULL,  
  is_correct BOOLEAN NOT NULL  
);
```

```
CREATE TABLE IF NOT EXISTS question_publications (  
  id BIGSERIAL PRIMARY KEY,  
  question_id BIGINT NOT NULL REFERENCES questions(id),
```

```
live_version INT NOT NULL,  
exam_code TEXT NOT NULL,  
tenant_id UUID NOT NULL,  
published_at TIMESTAMPTZ NOT NULL DEFAULT now(),  
UNIQUE (question_id, tenant_id, exam_code)  
);  
""")
```

```
write(f"{ROOT}/sql/delivery_ddl.sql", """)  
CREATE TABLE IF NOT EXISTS quiz_sessions (  
id UUID PRIMARY KEY,  
user_id UUID NOT NULL,  
tenant_id UUID NOT NULL,  
mode TEXT NOT NULL CHECK (mode IN ('tutor','exam')),  
adaptive BOOLEAN NOT NULL DEFAULT TRUE,  
exam_code TEXT,  
started_at TIMESTAMPTZ NOT NULL DEFAULT now(),  
expires_at TIMESTAMPTZ,  
sealed_at TIMESTAMPTZ  
);
```

```
CREATE TABLE IF NOT EXISTS quiz_items (  
id BIGSERIAL PRIMARY KEY,  
quiz_id UUID NOT NULL,  
question_id BIGINT NOT NULL,  
version INT NOT NULL,  
position INT NOT NULL,  
served_at TIMESTAMPTZ NOT NULL DEFAULT now(),  
UNIQUE(quiz_id, position)  
);
```

```
CREATE TABLE IF NOT EXISTS user_responses (  
id BIGSERIAL PRIMARY KEY,  
quiz_id UUID NOT NULL,  
user_id UUID NOT NULL,  
question_id BIGINT NOT NULL,  
version INT NOT NULL,  
option_label CHAR(1) NOT NULL,  
is_correct BOOLEAN NOT NULL,  
time_taken_ms INT,  
created_at TIMESTAMPTZ NOT NULL DEFAULT now()  
);  
""")
```

```

write(f"{ROOT}/sql/analytics_ddl.sql", """
CREATE TABLE IF NOT EXISTS item_calibration (
question_id BIGINT NOT NULL,
version INT NOT NULL,
model TEXT NOT NULL,
a FLOAT, b FLOAT, c FLOAT,
se_a FLOAT, se_b FLOAT, se_c FLOAT,
n_respondents INT,
fitted_at TIMESTAMPTZ NOT NULL DEFAULT now(),
PRIMARY KEY (question_id, version, model)
);
""")

```

```

write(f"{ROOT}/sql/indexes.sql", """
CREATE INDEX IF NOT EXISTS idx_qv_state ON question_versions(state);
CREATE INDEX IF NOT EXISTS idx_qv_topic ON question_versions(topic_id);
CREATE INDEX IF NOT EXISTS idx_qv_tags_gin ON question_versions USING GIN
(tags jsonb_path_ops);
CREATE INDEX IF NOT EXISTS idx_resp_user ON user_responses(user_id);
CREATE INDEX IF NOT EXISTS idx_resp_question ON user_responses(question_id,
version);
CREATE INDEX IF NOT EXISTS idx_quiz_items_qid_pos ON quiz_items(quiz_id,
position);
""")

```

```

write(f"{ROOT}/sql/auth_ddl.sql", """
-- Simple users with roles array for demo. Replace with IdP in prod.
CREATE TABLE IF NOT EXISTS users (
id UUID PRIMARY KEY,
email TEXT UNIQUE,
display_name TEXT,
roles TEXT[] NOT NULL DEFAULT ARRAY['student']
);
""")

```

----- Alembic -----

```

write(f"{ROOT}/alembic/alembic.ini", """
[alembic]
script_location = alembic
sqlalchemy.url = postgresql+psycpg2://qbank:qbank@localhost:5432/qbank
""")

```

```

write(f"{ROOT}/alembic/env.py", """

```



```

from logging.config import fileConfig
from sqlalchemy import engine_from_config, pool
from alembic import context
config = context.config
fileConfig(config.config_file_name)
target_metadata = None

def run_migrations_offline():
    context.configure(url=config.get_main_option("sqlalchemy.url"),
        literal_binds=True)
    with context.begin_transaction():
        context.run_migrations()

def run_migrations_online():
    connectable = engine_from_config(
        config.get_section(config.config_ini_section),
        prefix='sqlalchemy.',
        poolclass=pool.NullPool,
    )
    with connectable.connect() as connection:
        context.configure(connection=connection, target_metadata=target_metadata)
        with context.begin_transaction():
            context.run_migrations()

if context.is_offline_mode():
    run_migrations_offline()
else:
    run_migrations_online()
"""

write(f"{ROOT}/alembic/versions/001_sql_bootstrap.py", """
from alembic import op
import pathlib

revision = '001_sql_bootstrap'
down_revision = None
branch_labels = None
depends_on = None

def _run_sql(filename):
    root = pathlib.Path(file).resolve().parents[2] / 'sql' / filename
    with open(root, 'r', encoding='utf-8') as f:
        sql = fxread()
    op.execute(sql)

```

```
def upgrade():
    _run_sql('content_ddl.sql')
    _run_sql('delivery_ddl.sql')
    _run_sql('analytics_ddl.sql')
    _run_sql('indexes.sql')
    _run_sql('auth_ddl.sql')
```

```
def downgrade():
    pass
    """
```

```
----- Backend requirements -----
write(f"{ROOT}/qbank-backend/requirements.txt", """
fastapi==0.115.0
uvicorn==0.30.6
pydantic==2.8.2
python-dotenv==1.0.1
redis==5.0.8
kafka-python==2.0.2
psycopg2-binary==2.9.9
SQLAlchemy==2.0.32
PyJWT==2.9.0
""")
```

```
----- Backend core config -----
write(f"{ROOT}/qbank-backend/app/core/config.py", """
import os
from dotenv import load_dotenv
load_dotenv()
```

```
DATABASE_URL = os.getenv("DATABASE_URL", "postgresql+psycopg2://
qbank:qbank@localhost:5432/qbank")
REDIS_URL = os.getenv("REDIS_URL", "redis://localhost:6379/0")
KAFKA_BOOTSTRAP = os.getenv("KAFKA_BOOTSTRAP", "localhost:9092")
KAFKA_TOPIC_EVENTS = os.getenv("KAFKA_TOPIC_EVENTS", "events.qbank")
TENANT_ID = os.getenv("APP_TENANT_ID",
"00000000-0000-0000-0000-000000000001")
APP_SECRET = os.getenv("APP_SECRET", "dev-secret-change-me")
MAX_DAILY_EXPOSURES = int(os.getenv("MAX_DAILY_EXPOSURES", "500"))
""")
```

```
write(f"{ROOT}/qbank-backend/app/core/database.py", """
from sqlalchemy import create_engine
```

```

from sqlalchemy.orm import sessionmaker
from app.core.config import DATABASE_URL

engine = create_engine(DATABASE_URL, future=True, pool_pre_ping=True)
SessionLocal = sessionmaker(bind=engine, autocommit=False, autoflush=False,
future=True)

def get_db():
    db = SessionLocal()
    try:
        yield db
    finally:
        db.close()
    """

write(f"{ROOT}/qbank-backend/app/core/cache.py", """
import redis
from datetime import datetime
from app.core.config import REDIS_URL, MAX_DAILY_EXPOSURES

redis_client = redis.Redis.from_url(REDIS_URL, decode_responses=True)

def exposure_key(question_id: int, version: int) -> str:
    day = datetime.utcnow().strftime("%Y%m%d")
    return f"exp:{day}:{question_id}:{version}"

def can_serve(question_id: int, version: int) -> bool:
    key = exposure_key(question_id, version)
    count = int(redis_client.get(key) or 0)
    return count < MAX_DAILY_EXPOSURES

def bump_exposure(question_id: int, version: int) -> None:
    key = exposure_key(question_id, version)
    pipe = redis_client.pipeline()
    pipe.incr(key, 1)
    pipe.expire(key, 86400) # 1 day
    pipe.execute()
    """

write(f"{ROOT}/qbank-backend/app/core/events.py", """
import json
from uuid import uuid4
from datetime import datetime, timezone
from kafka import KafkaProducer

```

```

from app.core.config import KAFKA_BOOTSTRAP, KAFKA_TOPIC_EVENTS

producer = KafkaProducer(
    bootstrap_servers=KAFKA_BOOTSTRAP,
    value_serializer=lambda v: json.dumps(v).encode("utf-8"),
    key_serializer=lambda k: (k or "").encode("utf-8"),
    linger_ms=50, acks='1'
)

def emit(event_type: str, payload: dict):
    env = {"event_id": str(uuid4()), "event_type": event_type, "timestamp":
datetime.now(timezone.utc).isoformat(), **payload}
    producer.send(KAFKA_TOPIC_EVENTS, key=env["event_id"], value=env
    "")

----- Auth -----
write(f"{ROOT}/qbank-backend/app/core/auth.py", """
from fastapi import Depends, HTTPException, status
from fastapi.security import HTTPBearer, HTTPAuthorizationCredentials
from pydantic import BaseModel
from typing import List
import jwt
from datetime import datetime, timedelta, timezone
from app.core.config import APP_SECRET

class TokenData(BaseModel):
    sub: str
    roles: List[str]

bearer = HTTPBearer()

def create_token(user_id: str, roles: List[str], ttl_minutes: int = 120) -> str:
    now = datetime.now(timezone.utc)
    payload = {"sub": user_id, "roles": roles, "iat": int(now.timestamp()), "exp":
int((now + timedelta(minutes=ttl_minutes)).timestamp())}
    return jwt.encode(payload, APP_SECRET, algorithm="HS256")

def get_current_user(creds: HTTPAuthorizationCredentials = Depends(bearer)) ->
TokenData:
    try:
        payload = jwt.decode(creds.credentials, APP_SECRET, algorithms=["HS256"])
        return TokenData(sub=payload["sub"], roles=payload.get("roles", []))
    except Exception:
        raise HTTPException(status_code=status.HTTP_401_UNAUTHORIZED,

```

```
detail="Invalid or expired token")
```

```
def require_roles(*required: str):  
def checker(user: TokenData = Depends(get_current_user)):  
    roles = set(user.roles)  
    if not roles.intersection(set(required)):  
        raise HTTPException(status_code=403, detail="Insufficient role")  
    return user  
return checker  
"""
```

----- ORM -----

```
write(f"{ROOT}/qbank-backend/app/models/orm.py", """  
from sqlalchemy.orm import DeclarativeBase, Mapped, mapped_column  
from sqlalchemy import BigInteger, Integer, String, Text, Boolean, ForeignKey,  
JSON
```

```
class Base(DeclarativeBase): pass
```

```
class Topic(Base):  
    tablename = "topics"  
    id: Mapped[int] = mapped_column(BigInteger, primary_key=True)  
    tenant_id: Mapped[str] = mapped_column(String)  
    parent_id: Mapped[int | None] = mapped_column(BigInteger,  
        ForeignKey("topics.id"), nullable=True)  
    name: Mapped[str] = mapped_column(String)  
    blueprint_code: Mapped[str | None] = mapped_column(String, nullable=True)
```

```
class Question(Base):  
    tablename = "questions"  
    id: Mapped[int] = mapped_column(BigInteger, primary_key=True)  
    tenant_id: Mapped[str] = mapped_column(String)  
    external_ref: Mapped[str | None] = mapped_column(String, nullable=True)  
    created_by: Mapped[str] = mapped_column(String)  
    is_deleted: Mapped[bool] = mapped_column(Boolean, default=False)
```

```
class QuestionVersion(Base):  
    tablename = "question_versions"  
    id: Mapped[int] = mapped_column(BigInteger, primary_key=True)  
    question_id: Mapped[int] = mapped_column(BigInteger,  
        ForeignKey("questions.id"))  
    version: Mapped[int] = mapped_column(Integer)  
    state: Mapped[str] = mapped_column(String)  
    stem_md: Mapped[str] = mapped_column(Text)
```

```
lead_in: Mapped[str] = mapped_column(Text)
rationale_md: Mapped[str] = mapped_column(Text)
difficulty_label: Mapped[str | None] = mapped_column(String, nullable=True)
bloom_level: Mapped[int | None] = mapped_column(Integer, nullable=True)
topic_id: Mapped[int | None] = mapped_column(BigInteger,
ForeignKey("topics.id"), nullable=True)
tags: Mapped[dict] = mapped_column(JSON)
assets: Mapped[list] = mapped_column(JSON)
references: Mapped[list] = mapped_column(JSON)
```

```
class QuestionOption(Base):
    tablename = "question_options"
    id: Mapped[int] = mapped_column(BigInteger, primary_key=True)
    question_version_id: Mapped[int] = mapped_column(BigInteger,
ForeignKey("question_versions.id"))
    option_label: Mapped[str] = mapped_column(String(1))
    option_text_md: Mapped[str] = mapped_column(Text)
    is_correct: Mapped[bool] = mapped_column(Boolean)
```

```
class QuestionPublication(Base):
    tablename = "question_publications"
    id: Mapped[int] = mapped_column(BigInteger, primary_key=True)
    question_id: Mapped[int] = mapped_column(BigInteger,
ForeignKey("questions.id"))
    live_version: Mapped[int] = mapped_column(Integer)
    exam_code: Mapped[str] = mapped_column(String)
    tenant_id: Mapped[str] = mapped_column(String)
```

```
class QuizSession(Base):
    tablename = "quiz_sessions"
    id: Mapped[str] = mapped_column(String, primary_key=True)
    user_id: Mapped[str] = mapped_column(String)
    tenant_id: Mapped[str] = mapped_column(String)
    mode: Mapped[str] = mapped_column(String)
    adaptive: Mapped[bool] = mapped_column(Boolean, default=True)
    exam_code: Mapped[str | None] = mapped_column(String, nullable=True)
```

```
class QuizItem(Base):
    tablename = "quiz_items"
    id: Mapped[int] = mapped_column(BigInteger, primary_key=True)
    quiz_id: Mapped[str] = mapped_column(String)
    question_id: Mapped[int] = mapped_column(BigInteger)
    version: Mapped[int] = mapped_column(Integer)
    position: Mapped[int] = mapped_column(Integer)
```

```

class UserResponse(Base):
    tablename = "user_responses"
    id: Mapped[int] = mapped_column(BigInteger, primary_key=True)
    quiz_id: Mapped[str] = mapped_column(String)
    user_id: Mapped[str] = mapped_column(String)
    question_id: Mapped[int] = mapped_column(BigInteger)
    version: Mapped[int] = mapped_column(Integer)
    option_label: Mapped[str] = mapped_column(String(1))
    is_correct: Mapped[bool] = mapped_column(Boolean)
    time_taken_ms: Mapped[int | None] = mapped_column(Integer, nullable=True)

```

```

class ItemCalibration(Base):
    tablename = "item_calibration"
    question_id: Mapped[int] = mapped_column(BigInteger, primary_key=True)
    version: Mapped[int] = mapped_column(Integer, primary_key=True)
    model: Mapped[str] = mapped_column(String, primary_key=True)
    a: Mapped[float | None] = mapped_column()
    b: Mapped[float | None] = mapped_column()
    c: Mapped[float | None] = mapped_column()
    n_respondents: Mapped[int | None] = mapped_column(Integer)
    """

```

----- Services: adaptive with quotas/exposure -----

```

write(f"{ROOT}/qbank-backend/app/services/adaptive.py", """

```

```

import math, json

```

```

from typing import List, Dict, Optional, Tuple

```

```

from app.core.cache import can_serve

```

```

D = 1.7

```

```

def logistic(x: float) -> float: return 1.0 / (1.0 + math.exp(-x))

```

```

def prob_3pl(theta: float, a: float, b: float, c: float) -> float:

```

```

    return c + (1.0 - c) * logistic(D * a * (theta - b))

```

```

def fisher_info_3pl(theta: float, a: float, b: float, c: float) -> float:

```

```

    P = prob_3pl(theta, a, b, c); Q = 1.0 - P

```

```

    if P<=0 or Q<=0 or (1.0-c)<=0: return 0.0

```

```

    return (D2)*(a2)(Q/P)((P-c)/(1.0-c))**2

```

```

def quota_favors(candidate_topic_id: int, remaining_quota: Dict[str,int]) -> bool:

```

```

    # remaining_quota keyed by str(topic_id)

```

```

    return remaining_quota.get(str(candidate_topic_id), 0) > 0

```

```

def select_with_constraints(candidates: List[Dict], theta: float, remaining_quota:

```

```

    Dict[str,int]) -> Optional[Dict]:

```

```

best, best_l = None, -1.0
for it in candidates:
    if not can_serve(it["question_id"], it["version"]): # exposure gate
        continue
    # prefer items from topics with remaining quotas
    topic_bonus = 0.5 if quota_favors(it.get("topic_id"), remaining_quota) else 0.0
    a = it.get("a", 1.0); b = it.get("b", 0.0); c = it.get("c", 0.2)
    l = fisher_info_3pl(theta, a, b, c) + topic_bonus
    if l > best_l:
        best_l, best = l, it
    return best
"""

```

----- API: auth, authoring, quizzes -----

```

write(f"{ROOT}/qbank-backend/app/api/auth.py", """
from fastapi import APIRouter
from pydantic import BaseModel
from typing import List
from app.core.auth import create_token

```

```

router = APIRouter()

```

```

class MockLogin(BaseModel):
    user_id: str
    roles: List[str]

```

```

@router.post("/mock-login")
def mock_login(payload: MockLogin):
    token = create_token(payload.user_id, payload.roles)
    return {"access_token": token, "token_type": "bearer", "roles": payload.roles}
"""

```

```

write(f"{ROOT}/qbank-backend/app/api/author.py", """
from fastapi import APIRouter, Depends, HTTPException
from pydantic import BaseModel, constr
from typing import List, Optional
from sqlalchemy.orm import Session
from sqlalchemy import select, func

```

```

from app.core.database import get_db
from app.core.config import TENANT_ID
from app.core.auth import require_roles, TokenData
from app.models.orm import Topic, Question, QuestionVersion, QuestionOption,
QuestionPublication

```



```
router = APIRouter()
```

```
class OptionIn(BaseModel):  
    label: constr(min_length=1, max_length=1)  
    text_md: str  
    is_correct: bool
```

```
class QuestionCreate(BaseModel):  
    external_ref: Optional[str] = None  
    topic_name: str  
    exam_code: str = "DEMO-EXAM"  
    stem_md: str  
    lead_in: str  
    rationale_md: str  
    difficulty_label: Optional[str] = "medium"  
    options: List[OptionIn]
```

```
@router.post("/questions",  
dependencies=[Depends(require_roles("author","admin"))])  
def create_question(payload: QuestionCreate, user: TokenData =  
Depends(require_roles("author","admin")), db: Session = Depends(get_db)):  
    t = db.scalar(select(Topic)×where(Topic×name == payload.topic_name))  
    if not t:  
        t = Topic(tenant_id=TENANT_ID, parent_id=None, name=payload×topic_name,  
blueprint_code=None)  
        db.add(t); db.flush()
```

python  
Always show details

Copy

```
q = Question(tenant_id=TENANT_ID, external_ref=payload.external_ref,  
created_by=user.sub, is_deleted=False)  
db.add(q); db.flush()
```

```
next_v = (db.scalar(select(func×coalesce(func×max(QuestionVersion×version),  
0))×where(QuestionVersion×question_id == q.id)) or 0) + 1  
qv = QuestionVersion(  
    question_id=q×id, version=next_v, state="published",  
    stem_md=payload.stem_md, lead_in=payload.lead_in,  
    rationale_md=payload.rationale_md,  
    difficulty_label=payload×difficulty_label, topic_id=txid, tags={}, assets=[],  
    references=[])
```

```

)
db.add(qv); db.flush()

for o in payload.options:
    db.add(QuestionOption(question_version_id=qv×id,
        option_label=o.label.upper(), option_text_md=o.text_md,
        is_correct=o.is_correct))

db.add(QuestionPublication(question_id=qxid, live_version=next_v,
    exam_code=payload.exam_code, tenant_id=TENANT_ID))
db.commit()
return {"question_id": q.id, "version": next_v, "topic_id": t.id}
@router.post("/publish/{question_id}",
    dependencies=[Depends(require_roles("publisher","admin"))])
def publish(question_id: int, exam_code: str = "DEMO-EXAM", db: Session =
    Depends(get_db)):
    qv = db×scalar(select(QuestionVersion)×where(QuestionVersion×question_id ==
        question_id).order_by(QuestionVersion.version.desc()))
    if not qv:
        raise HTTPException(404, "Question not found")
    pub = QuestionPublication(question_id=question_id, live_version=qv×version,
        exam_code=exam_code, tenant_id=TENANT_ID)
    db.add(pub); db.commit()
    return {"published": True, "question_id": question_id, "version": qv.version,
        "exam_code": exam_code}
"""

```

```

write(f"{ROOT}/qbank-backend/app/api/quizzes.py", """
from fastapi import APIRouter, HTTPException, Depends
from pydantic import BaseModel, Field, constr
from typing import List, Optional, Literal, Dict
from uuid import uuid4
from datetime import datetime, timedelta
import json
from sqlalchemy.orm import Session
from sqlalchemy import select
from app.core.cache import redis_client, bump_exposure
from app.core.events import emit
from app.core.database import get_db
from app.core.auth import require_roles, TokenData
from app.models.orm import QuestionVersion, QuestionOption,
    QuestionPublication, ItemCalibration
from app.services.adaptive import select_with_constraints

```

```
router = APIRouter()
```

```
class QuizFilters(BaseModel):  
    topics: Optional[List[str]] = None  
    difficulty: Optional[List[Literal["easy","medium","hard"]]] = None  
    num_questions: int = Field(ge=1, le=120, default=40)  
    mode: Literal["tutor","exam"] = "tutor"  
    exam_code: Optional[str] = "DEMO-EXAM"
```

```
class QuizCreate(BaseModel):  
    tenant_id: constr(min_length=8)  
    filters: QuizFilters  
    adaptive: bool = True  
    blueprint_quota: Optional[Dict[str,int]] = None # topic_name -> count
```

```
class QuizCreated(BaseModel):  
    quiz_id: str  
    question_ids: List[int]  
    expires_at: datetime  
    mode: Literal["tutor","exam"]
```

```
class NextQuestion(BaseModel):  
    question_id: int  
    version: int  
    payload: dict
```

```
class AnswerSubmit(BaseModel):  
    question_id: int  
    selected: constr(min_length=1, max_length=1)  
    time_taken_ms: Optional[int] = 0  
    client_latency_ms: Optional[int] = 0
```

```
class AnswerResult(BaseModel):  
    correct: bool  
    correct_option: constr(min_length=1, max_length=1)  
    explanation: dict  
    difficulty: float
```

```
def _rk(qid: str, suf: str) -> str: return f"quiz:{qid}:{suf}"
```

```
@router.post("", response_model=QuizCreated, status_code=201,  
dependencies=[Depends(require_roles("student","admin"))])  
def create_quiz(payload: QuizCreate, user: TokenData =
```

```
Depends(require_roles("student","admin")), db: Session = Depends(get_db)):
quiz_id = str(uuid4()); mode = payload.filters.mode
expires_at = datetime.utcnow() + timedelta(hours=2)
```

```
stmt = select(QuestionPublication, QuestionVersion).join(
    QuestionVersion,
    (QuestionVersion.question_id == QuestionPublication.question_id) &
    (QuestionVersion.version == QuestionPublication.live_version)
).where(QuestionPublication.exam_code == (payload.filters.exam_code or "DEMO-EXAM"), QuestionVersion.state == "published")
rows = db.execute(stmt).all()
versions = [r[1] for r in rows]
```

```
cache versions and planned quotas
vcache = [{"q": v.question_id, "v": v.version, "t": v.topic_id, "d": v.difficulty_label}
for v in versions]
redis_client.set(_rk(quiz_id, "versions"), json.dumps(vcache), ex=7200)
redis_client.set(_rk(quiz_id, "cursor"), 0, ex=7200)
redis_client.set(_rk(quiz_id, "mode"), mode, ex=7200)
redis_client.set(_rk(quiz_id, "user"), user.sub, ex=7200)
```

```
plan quotas by topic_id mapping from names, simple proportional fallback
quotas = {}
if payload.blueprint_quota:
    # Convert topic names to ids (best effort)
    # For simplicity we store by topic_id string; caller may pass ids already
    # Here we just pass through names as keys; selection will compare string(topic_id)
    for k,v in payload.blueprint_quota.items(): quotas[str(k)] = int(v)
else:
    # default: even distribution across topics present
    topic_counts = {}
    for v in vcache: topic_counts[str(v["t"])] = topic_counts.get(str(v["t"]),0)+1
    topics = list(topic_counts.keys())
    if topics:
        per = max(1, payload.filters.num_questions // max(1,len(topics)))
        for t in topics: quotas[str(t)] = per
    redis_client.set(_rk(quiz_id,"quota_remaining"), json.dumps(quotas), ex=7200)
```

```
emit("quiz_started", {"quiz_id": quiz_id, "user_id": user.sub, "tenant_id":
payload.tenant_id, "mode": mode, "filters": payload.filters.model_dump(), "quota":
quotas})
qids = list({v["q"] for v in vcache}[:payload.filters.num_questions])
return QuizCreated(quiz_id=quiz_id, question_ids=qids, expires_at=expires_at,
mode=mode)
```

```

@router.get("/{quiz_id}/next", response_model=NextQuestion,
dependencies=[Depends(require_roles("student","admin"))])
def next_question(quiz_id: str, db: Session = Depends(get_db)):
    raw = redis_client.get(_rk(quiz_id,"versions"))
    if not raw: raise HTTPException(404, "Quiz not found or expired")
    versions = json.loads(raw)

    curk = _rk(quiz_id, "cursor")
    cur = int(redis_client.get(curk) or 0)
    if cur >= len(versions): raise HTTPException(404, "No more questions")

    quota = json.loads(redis_client.get(_rk(quiz_id,"quota_remaining")) or "{}")

    build a candidate window and hydrate IRT
    window = versions[cur : min(cur+20, len(versions))]
    candidates = []
    for w in window:
        ic =
        db.scalar(select(ItemCalibration).where(ItemCalibration.question_id==w["q"],
        ItemCalibration.version==w["v"]).limit(1))
        if ic:
            a = ic.a or 1.0; b = ic.b or 0.0; c = ic.c or (0.2 if (ic.model or "3PL")=="3PL" else
            0.0)
        else:
            a,b,c = 1.0,0.0,0.2
        candidates.append({"question_id": w["q"], "version": w["v"], "topic_id": w["t"],
        "a": a, "b": b, "c": c})

    best = select_with_constraints(candidates, theta=0.0, remaining_quota=quota) or
    candidates[0]

    decrement quota for selected topic
    tkey = str(best.get("topic_id"))
    if tkey in quota and quota[tkey] > 0:
        quota[tkey] -= 1
        redis_client.set(_rk(quiz_id,"quota_remaining"), json.dumps(quota))
        redis_client.set(curk, cur+1)

    fetch payload
    qv =
    db.scalar(select(QuestionVersion).where(QuestionVersion.question_id==best["que
    stion_id"], QuestionVersion.version==best["version"]))
    if not qv: raise HTTPException(500, "Item not found")

```

```

opts =
db.execute(select(QuestionOption).where(QuestionOption.question_version_id==
qv.id)).scalars().all()
bump_exposure(best["question_id"], best["version"])

```

```

payload = {"stem_md": qv.stem_md, "lead_in": qv.lead_in, "options": [{"label":
o.option_label, "text": o.option_text_md} for o in opts]}
emit("question_served", {"quiz_id": quiz_id, "question_id": best["question_id"],
"version": best["version"]})
return NextQuestion(question_id=best["question_id"], version=best["version"],
payload=payload)

```

```

@router.post("/{quiz_id}/answers", response_model=AnswerResult,
dependencies=[Depends(require_roles("student","admin"))])
def submit_answer(quiz_id: str, payload: AnswerSubmit, db: Session =
Depends(get_db)):
qv =
db.scalar(select(QuestionVersion).where(QuestionVersion.question_id==payload.q
uestion_id).order_by(QuestionVersion.version.desc()))
if not qv: raise HTTPException(404, "Question not found")
opts =
db.execute(select(QuestionOption).where(QuestionOption.question_version_id==
qv.id)).scalars().all()
correct = next((o.option_label for o in opts if o.is_correct), None)
if not correct: raise HTTPException(500, "No correct option set")
ok = (payload.selected.upper() == correct)
emit("answer_submitted", {"quiz_id": quiz_id, "question_id": payload.question_id,
"version": qv.version, "selected": payload.selected, "is_correct": ok,
"time_taken_ms": payload.time_taken_ms or 0})
return AnswerResult(correct=ok, correct_option=correct,
explanation={"rationale_md": qv.rationale_md}, difficulty=0.5)
"""

```

----- app main -----

```

write(f"{ROOT}/qbank-backend/app/main.py", """
from fastapi import FastAPI
from fastapi.middleware.cors import CORSMiddleware
from app.api.quizzes import router as quizzes_router
from app.api.author import router as author_router
from app.api.auth import router as auth_router

app = FastAPI(title="QBank API v3", version="3.0.0")
app.add_middleware(CORSMiddleware, allow_origins=[""],
allow_credentials=True, allow_methods=["*"], allow_headers=["*"])

```

```
app.include_router(auth_router, prefix="/v1/auth", tags=["auth"])
app.include_router(quizzes_router, prefix="/v1/quizzes", tags=["quizzes"])
app.include_router(author_router, prefix="/v1/author", tags=["authoring"])
```

```
@app.get("/health")
def health(): return {"status": "ok"}
"""
```

```
----- dbt additions -----
write(f"{ROOT}/dbt-qbank/dbt_project.yml", """
name: "dbt_qbank"
version: "1.0.0"
config-version: 2
profile: "dbt_qbank"
```

```
model-paths: ["models"]
target-path: "target"
clean-targets: ["target"]
```

```
models:
dbt_qbank:
+materialized: view
marts:
+materialized: table
""")
```

```
write(f"{ROOT}/dbt-qbank/profiles.yml", """
dbt_qbank:
outputs:
dev:
type: postgres
host: localhost
user: qbank
password: qbank
port: 5432
dbname: qbank
schema: analytics
target: dev
""")
```

```
write(f"{ROOT}/dbt-qbank/models/marts/cohort_performance.sql", """
-- Cohort performance per tenant/day
with resp as (
```

```

select user_id, question_id, version, is_correct, created_at::date as d,
'default'::text as cohort
from public.user_responses
)
select
d as day,
cohort,
avg(case when is_correct then 1 else 0 end) as accuracy,
count(*) as attempts
from resp
group by 1,2
order by 1 desc
""")

```

----- Docker compose -----

```

write(f"{ROOT}/docker/docker-compose.yml", ""
services:
postgres:
image: postgres:16-alpine
environment:
POSTGRES_USER: qbank
POSTGRES_PASSWORD: qbank
POSTGRES_DB: qbank
ports: ["5432:5432"]
volumes: [ "pgdata:/var/lib/postgresql/data" ]

```

```

redis:
image: redis:7-alpine
ports: ["6379:6379"]

```

```

zookeeper:
image: confluentinc/cp-zookeeper:7.6.1
environment:
ZOOKEEPER_CLIENT_PORT: 2181
ZOOKEEPER_TICK_TIME: 2000
ports: ["2181:2181"]

```

```

kafka:
image: confluentinc/cp-kafka:7.6.1
depends_on: [zookeeper]
ports: ["9092:9092"]
environment:
KAFKA_BROKER_ID: 1
KAFKA_ZOOKEEPER_CONNECT: zookeeper:2181

```



```
KAFKA_ADVERTISED_LISTENERS: PLAINTEXT://kafka:9092,PLAINTEXT_HOST://localhost:9092
KAFKA_LISTENER_SECURITY_PROTOCOL_MAP:
PLAINTEXT:PLAINTEXT,PLAINTEXT_HOST:PLAINTEXT
KAFKA_INTER_BROKER_LISTENER_NAME: PLAINTEXT
KAFKA_OFFSETS_TOPIC_REPLICATION_FACTOR: 1
```

```
kafdrop:
image: obsidiandynamics/kafdrop:4.0.1
depends_on: [kafka]
ports: ["19000:9000"]
environment:
KAFKA_BROKERCONNECT: "kafka:9092"
```

volumes:

```
pgdata:
""")
```

----- Zip -----

```
zip_path = "/mnt/data/qbank_scaffold_v3.zip"
with zipfile.ZipFile(zip_path, "w", zipfile.ZIP_DEFLATED) as z:
for p in pathlib.Path(ROOT).rglob("*"):
z.write(p, p.relative_to(ROOT))
```

zip\_path

Always show details

Copy

Result

'/mnt/data/qbank\_scaffold\_v3.zip'

Delivered. I've pushed the next set of production-grade upgrades and packed them up for you:

Download the scaffold (v3)

What's new in v3:

JWT auth + RBAC: roles (author, publisher, student, admin) enforced on endpoints.

Alembic migration scripts (SQL-first) to bootstrap your DB.

Adaptive engine upgrades: per-item daily exposure cap (Redis) + blueprint quota targeting by topic.

Extended ORM incl. quiz sessions/items, responses, calibration.

dbt cohort analytics (tenant/day rollups).

How to try it fast:

Infra

```
cd docker && export APP_SECRET="change-me-please" && docker compose up -d
```

DB init

```
psql postgresql://qbank:qbank@localhost:5432/qbank -f ../sql/content_ddl.sql
psql postgresql://qbank:qbank@localhost:5432/qbank -f ../sql/delivery_ddl.sql
psql postgresql://qbank:qbank@localhost:5432/qbank -f ../sql/analytics_ddl.sql
psql postgresql://qbank:qbank@localhost:5432/qbank -f ../sql/indexes.sql
psql postgresql://qbank:qbank@localhost:5432/qbank -f ../sql/auth_ddl.sql
```

API

```
cd ../qbank-backend && python -m venv .venv && source .venv/bin/activate && pip
install -r requirements.txt && uvicorn app.main:app --reload
Open http://localhost:8000/docs
```

Get a token (mock)

```
POST /v1/auth/mock-login with {"user_id":"user-123","roles":
["author","publisher","student"]}
```

Create & publish an item (auth required)

```
POST /v1/author/questions → returns question_id
POST /v1/author/publish/{question_id}
```

Start a quiz with quotas

```
POST /v1/quizzes (pass blueprint_quota to enforce per-topic counts)
Then call /v1/quizzes/{id}/next and /answers.
```

Want me to:

add a full SQLAlchemy migration path (Alembic autogenerate) and CI scripts,

implement item exposure control using Simpson–Hetter proportions,

wire a tiny Next.js student UI to take quizzes,

# QBank Scaffold v4:

# - Alembic autogenerate path + CI workflows

```
# - Sympson–Hetter exposure control
# - Next.js student UI to take quizzes
# - Looker/Metabase models over dbt marts
#
# Outputs /mnt/data/qbank_scaffold_v4.zip
```

```
import os, zipfile, pathlib, textwrap, json
```

```
ROOT = "/mnt/data/qbank_scaffold_v4"
os.makedirs(ROOT, exist_ok=True)
```

```
def write(path, content):
    os.makedirs(os.path.dirname(path), exist_ok=True)
    with open(path, "w", encoding="utf-8") as f:
        f.write(textwrap.dedent(content))
```

```
# ----- README -----
write(f"{ROOT}/README.md", """
# QBank Backend Scaffold (v4)
```

Adds:

- **Alembic autogenerate** wired to SQLAlchemy models (`Base.metadata``), with example generated migration.
- **CI scripts** (GitHub Actions): backend pipeline (lint, alembic upgrade head with Postgres service), student UI build.
- **Sympson–Hetter exposure control**: per-item `sh_p`` acceptance sampling to limit overexposure.
- **Student UI (Next.js)**: create quiz, fetch next, submit answers.
- **Looker** (LookML) + **Metabase** (API import scripts) models over dbt marts.

> This builds on v3. You can drop v4 into the same repo or start fresh.

## ## Quick Start

- 1) Infra (from v3 instructions) — bring up Postgres/Redis/Kafka.
- 2) DB init: run SQL in `/sql`` (includes `item_exposure_control.sql``).
- 3) Backend: `cd qbank-backend && pip install -r requirements.txt && uvicorn app.main:app --reload``.
- 4) Student UI: `cd student-ui && npm install && npm run dev`` → `http://localhost:3000`
- 5) CI (GitHub Actions): copy `.github/workflows/*`` into your repo.

## ## Sympson–Hetter

- Table: `item_exposure_control(question_id, version, sh_p)``
- Selector samples `u~U(0,1)`` and serves item only if `u <= sh_p``; otherwise picks

next best candidate.

- Default `sh\_p = 1.0` if no row exists.

```
""")
```

```
# ----- SQL: add item_exposure_control -----
```

```
write(f"{ROOT}/sql/item_exposure_control.sql", """
```

```
CREATE TABLE IF NOT EXISTS item_exposure_control (
```

```
    question_id BIGINT NOT NULL,
```

```
    version INT NOT NULL,
```

```
    sh_p FLOAT NOT NULL DEFAULT 1.0, -- Simpson-Hetter acceptance probability
```

```
    updated_at TIMESTAMPTZ NOT NULL DEFAULT now(),
```

```
    PRIMARY KEY (question_id, version)
```

```
);
```

```
""")
```

```
# ----- Backend requirements (ensure same as v3 + extras) -----
```

```
write(f"{ROOT}/qbank-backend/requirements.txt", """
```

```
fastapi==0.115.0
```

```
uvicorn==0.30.6
```

```
pydantic==2.8.2
```

```
python-dotenv==1.0.1
```

```
redis==5.0.8
```

```
kafka-python==2.0.2
```

```
psycopg2-binary==2.9.9
```

```
SQLAlchemy==2.0.32
```

```
PyJWT==2.9.0
```

```
alembic==1.13.2
```

```
black==24.8.0
```

```
flake8==7.1.1
```

```
""")
```

```
# ----- Backend: Alembic configured for autogenerate -----
```

```
write(f"{ROOT}/qbank-backend/alembic.ini", """
```

```
[alembic]
```

```
script_location = alembic
```

```
[loggers]
```

```
keys = root,sqlalchemy,alembic
```

```
[handlers]
```

```
keys = console
```

```
[formatters]
```

```
keys = generic
```

```
[logger_root]
level = WARN
handlers = console
```

```
[logger_sqlalchemy]
level = WARN
handlers = console
qualname = sqlalchemy.engine
```

```
[logger_alembic]
level = INFO
handlers = console
qualname = alembic
```

```
[handler_console]
class = StreamHandler
args = (sys.stderr,)
level = NOTSET
formatter = generic
```

```
[formatter_generic]
format = %(levelname)-5.5s [%(name)s] %(message)s
"""
```

```
write(f"{ROOT}/qbank-backend/alembic/env.py", """
from logging.config import fileConfig
from sqlalchemy import engine_from_config, pool
from alembic import context
import os
```

```
# Import your models' metadata
from app.models.orm import Base
from app.core.config import DATABASE_URL
```

```
config = context.config
fileConfig(config.config_file_name)
target_metadata = Base.metadata
```

```
def run_migrations_offline():
    url = DATABASE_URL
    context.configure(
        url=url,
        target_metadata=target_metadata,
```

```

        literal_binds=True,
        compare_type=True,
        compare_server_default=True,
    )
    with context.begin_transaction():
        context.run_migrations()

def run_migrations_online():
    configuration = config.get_section(config.config_ini_section) or {}
    configuration["sqlalchemy.url"] = DATABASE_URL

    connectable = engine_from_config(
        configuration,
        prefix="sqlalchemy.",
        poolclass=pool.NullPool,
    )
    with connectable.connect() as connection:
        context.configure(
            connection=connection,
            target_metadata=target_metadata,
            compare_type=True,
            compare_server_default=True,
        )
        with context.begin_transaction():
            context.run_migrations()

if context.is_offline_mode():
    run_migrations_offline()
else:
    run_migrations_online()
"""

write(f"{ROOT}/qbank-backend/alembic/script.py.mako", """
<%text># Auto-generated Alembic migration script.</%text>
revision = '${up_revision}'
down_revision = ${repr(down_revision)}
branch_labels = ${repr(branch_labels)}
depends_on = ${repr(depends_on)}

from alembic import op
import sqlalchemy as sa

def upgrade():
    ${upgrades if upgrades else "pass"}

```

```

def downgrade():
    """
    ${downgrades if downgrades else "pass"}
    """

# Example pre-generated revision adding item_exposure_control (matches SQL
file)
write(f"{ROOT}/qbank-backend/alembic/versions/002_item_exposure_control.py",
"""
revision = '002_item_exposure_control'
down_revision = None
branch_labels = None
depends_on = None

from alembic import op
import sqlalchemy as sa

def upgrade():
    op.create_table(
        'item_exposure_control',
        sa.Column('question_id', sa.BigInteger(), nullable=False),
        sa.Column('version', sa.Integer(), nullable=False),
        sa.Column('sh_p', sa.Float(), nullable=False, server_default=sa.text('1.0')),
        sa.Column('updated_at', sa.TIMESTAMP(timezone=True),
server_default=sa.text('now()'), nullable=False),
        sa.PrimaryKeyConstraint('question_id','version')
    )

def downgrade():
    op.drop_table('item_exposure_control')
    """

# ----- Backend: ORM (includes ItemExposureControl) -----
write(f"{ROOT}/qbank-backend/app/models/orm.py", """
from sqlalchemy.orm import DeclarativeBase, Mapped, mapped_column
from sqlalchemy import BigInteger, Integer, String, Text, Boolean, ForeignKey,
JSON, Float

class Base(DeclarativeBase): pass

class Topic(Base):
    __tablename__ = "topics"
    id: Mapped[int] = mapped_column(BigInteger, primary_key=True)
    tenant_id: Mapped[str] = mapped_column(String)

```

```
parent_id: Mapped[int | None] = mapped_column(BigInteger,
ForeignKey("topics.id"), nullable=True)
name: Mapped[str] = mapped_column(String)
blueprint_code: Mapped[str | None] = mapped_column(String, nullable=True)
```

```
class Question(Base):
```

```
    __tablename__ = "questions"
    id: Mapped[int] = mapped_column(BigInteger, primary_key=True)
    tenant_id: Mapped[str] = mapped_column(String)
    external_ref: Mapped[str | None] = mapped_column(String, nullable=True)
    created_by: Mapped[str] = mapped_column(String)
    is_deleted: Mapped[bool] = mapped_column(Boolean, default=False)
```

```
class QuestionVersion(Base):
```

```
    __tablename__ = "question_versions"
    id: Mapped[int] = mapped_column(BigInteger, primary_key=True)
    question_id: Mapped[int] = mapped_column(BigInteger,
ForeignKey("questions.id"))
    version: Mapped[int] = mapped_column(Integer)
    state: Mapped[str] = mapped_column(String)
    stem_md: Mapped[str] = mapped_column(Text)
    lead_in: Mapped[str] = mapped_column(Text)
    rationale_md: Mapped[str] = mapped_column(Text)
    difficulty_label: Mapped[str | None] = mapped_column(String, nullable=True)
    bloom_level: Mapped[int | None] = mapped_column(Integer, nullable=True)
    topic_id: Mapped[int | None] = mapped_column(BigInteger,
ForeignKey("topics.id"), nullable=True)
    tags: Mapped[dict] = mapped_column(JSON)
    assets: Mapped[list] = mapped_column(JSON)
    references: Mapped[list] = mapped_column(JSON)
```

```
class QuestionOption(Base):
```

```
    __tablename__ = "question_options"
    id: Mapped[int] = mapped_column(BigInteger, primary_key=True)
    question_version_id: Mapped[int] = mapped_column(BigInteger,
ForeignKey("question_versions.id"))
    option_label: Mapped[str] = mapped_column(String(1))
    option_text_md: Mapped[str] = mapped_column(Text)
    is_correct: Mapped[bool] = mapped_column(Boolean)
```

```
class QuestionPublication(Base):
```

```
    __tablename__ = "question_publications"
    id: Mapped[int] = mapped_column(BigInteger, primary_key=True)
    question_id: Mapped[int] = mapped_column(BigInteger,
```



```
ForeignKey("questions.id"))
```

```
live_version: Mapped[int] = mapped_column(Integer)
```

```
exam_code: Mapped[str] = mapped_column(String)
```

```
tenant_id: Mapped[str] = mapped_column(String)
```

```
class QuizSession(Base):
```

```
    __tablename__ = "quiz_sessions"
```

```
    id: Mapped[str] = mapped_column(String, primary_key=True)
```

```
    user_id: Mapped[str] = mapped_column(String)
```

```
    tenant_id: Mapped[str] = mapped_column(String)
```

```
    mode: Mapped[str] = mapped_column(String)
```

```
    adaptive: Mapped[bool] = mapped_column(Boolean, default=True)
```

```
    exam_code: Mapped[str | None] = mapped_column(String, nullable=True)
```

```
class QuizItem(Base):
```

```
    __tablename__ = "quiz_items"
```

```
    id: Mapped[int] = mapped_column(BigInteger, primary_key=True)
```

```
    quiz_id: Mapped[str] = mapped_column(String)
```

```
    question_id: Mapped[int] = mapped_column(BigInteger)
```

```
    version: Mapped[int] = mapped_column(Integer)
```

```
    position: Mapped[int] = mapped_column(Integer)
```

```
class UserResponse(Base):
```

```
    __tablename__ = "user_responses"
```

```
    id: Mapped[int] = mapped_column(BigInteger, primary_key=True)
```

```
    quiz_id: Mapped[str] = mapped_column(String)
```

```
    user_id: Mapped[str] = mapped_column(String)
```

```
    question_id: Mapped[int] = mapped_column(BigInteger)
```

```
    version: Mapped[int] = mapped_column(Integer)
```

```
    option_label: Mapped[str] = mapped_column(String(1))
```

```
    is_correct: Mapped[bool] = mapped_column(Boolean)
```

```
    time_taken_ms: Mapped[int | None] = mapped_column(Integer, nullable=True)
```

```
class ItemCalibration(Base):
```

```
    __tablename__ = "item_calibration"
```

```
    question_id: Mapped[int] = mapped_column(BigInteger, primary_key=True)
```

```
    version: Mapped[int] = mapped_column(Integer, primary_key=True)
```

```
    model: Mapped[str] = mapped_column(String, primary_key=True)
```

```
    a: Mapped[float | None] = mapped_column(Float)
```

```
    b: Mapped[float | None] = mapped_column(Float)
```

```
    c: Mapped[float | None] = mapped_column(Float)
```

```
    n_respondents: Mapped[int | None] = mapped_column(Integer)
```

```
class ItemExposureControl(Base):
```

```

__tablename__ = "item_exposure_control"
question_id: Mapped[int] = mapped_column(BigInteger, primary_key=True)
version: Mapped[int] = mapped_column(Integer, primary_key=True)
sh_p: Mapped[float] = mapped_column(Float) # 0..1
"""

# ----- Backend: Adaptive selector w/ Sympton-Hetter -----
write(f"{ROOT}/qbank-backend/app/services/adaptive.py", """
import math, random
from typing import List, Dict, Optional
D = 1.7

def logistic(x: float) -> float: return 1.0 / (1.0 + math.exp(-x))
def prob_3pl(theta: float, a: float, b: float, c: float) -> float:
    return c + (1.0 - c) * logistic(D * a * (theta - b))
def fisher_info_3pl(theta: float, a: float, b: float, c: float) -> float:
    P = prob_3pl(theta, a, b, c); Q = 1.0 - P
    if P<=0 or Q<=0 or (1.0-c)<=0: return 0.0
    return (D**2)*(a**2)*(Q/P)*((P-c)/(1.0-c))**2

def select_with_SH_and_quota(candidates: List[Dict], theta: float,
remaining_quota: Dict[str,int]) -> Optional[Dict]:
    # Rank by Fisher information, then attempt SH acceptance; if rejected, try next.
    scored = []
    for it in candidates:
        a = it.get("a",1.0); b = it.get("b",0.0); c = it.get("c",0.2)
        I = fisher_info_3pl(theta, a, b, c)
        # slight boost if topic quota remains
        if remaining_quota.get(str(it.get("topic_id")),0) > 0:
            I += 0.25
        scored.append((I, it))
    scored=sorted(key=lambda x: x[0], reverse=True)

    for _, it in scored:
        sh_p = it.get("sh_p", 1.0)
        if random.random() <= max(0.0, min(1.0, sh_p)):
            return it # accepted by SH
    return scored[0][1] if scored else None # if all rejected, serve top (or None)
""")

# ----- Backend: API changes to supply sh_p -----
write(f"{ROOT}/qbank-backend/app/api/quizzes.py", """
from fastapi import APIRouter, HTTPException, Depends
from pydantic import BaseModel, Field, constr

```

```

from typing import List, Optional, Literal, Dict
from uuid import uuid4
from datetime import datetime, timedelta
import json
from sqlalchemy.orm import Session
from sqlalchemy import select
from app.core.cache import redis_client, bump_exposure
from app.core.events import emit
from app.core.database import get_db
from app.core.auth import require_roles, TokenData
from app.models.orm import QuestionVersion, QuestionOption,
QuestionPublication, ItemCalibration, ItemExposureControl
from app.services.adaptive import select_with_SH_and_quota

router = APIRouter()

class QuizFilters(BaseModel):
    topics: Optional[List[str]] = None
    difficulty: Optional[List[Literal["easy","medium","hard"]]] = None
    num_questions: int = Field(ge=1, le=120, default=40)
    mode: Literal["tutor","exam"] = "tutor"
    exam_code: Optional[str] = "DEMO-EXAM"

class QuizCreate(BaseModel):
    tenant_id: constr(min_length=8)
    filters: QuizFilters
    adaptive: bool = True
    blueprint_quota: Optional[Dict[str,int]] = None # topic_id (string) -> count

class QuizCreated(BaseModel):
    quiz_id: str
    question_ids: List[int]
    expires_at: datetime
    mode: Literal["tutor","exam"]

class NextQuestion(BaseModel):
    question_id: int
    version: int
    payload: dict

class AnswerSubmit(BaseModel):
    question_id: int
    selected: constr(min_length=1, max_length=1)
    time_taken_ms: Optional[int] = 0

```

client\_latency\_ms: Optional[int] = 0

class AnswerResult(BaseModel):

correct: bool

correct\_option: constr(min\_length=1, max\_length=1)

explanation: dict

difficulty: float

def \_rk(qid: str, suf: str) -> str: return f"quiz:{qid}:{suf}"

@router.post("", response\_model=QuizCreated, status\_code=201,

dependencies=[Depends(require\_roles("student","admin"))])

def create\_quiz(payload: QuizCreate, user: TokenData =

Depends(require\_roles("student","admin")), db: Session = Depends(get\_db)):

quiz\_id = str(uuid4()); mode = payload.filters.mode

expires\_at = datetime.utcnow() + timedelta(hours=2)

stmt = select(QuestionPublication, QuestionVersion).join(

QuestionVersion,

(QuestionVersion.question\_id == QuestionPublication.question\_id) &

(QuestionVersion.version == QuestionPublication.live\_version)

).where(QuestionPublication.exam\_code == (payload.filters.exam\_code or  
"DEMO-EXAM"), QuestionVersion.state == "published")

rows = db.execute(stmt).all()

versions = [r[1] for r in rows]

# cache versions list with minimal fields

vcache = [{"q": v.question\_id, "v": v.version, "t": v.topic\_id, "d": v.difficulty\_label}

for v in versions]

redis\_client.set(\_rk(quiz\_id, "versions"), json.dumps(vcache), ex=7200)

redis\_client.set(\_rk(quiz\_id, "cursor"), 0, ex=7200)

redis\_client.set(\_rk(quiz\_id, "mode"), mode, ex=7200)

redis\_client.set(\_rk(quiz\_id, "user"), user.sub, ex=7200)

# quotas (topic\_id string -> count)

quotas = payload.blueprint\_quota or {}

redis\_client.set(\_rk(quiz\_id,"quota\_remaining"), json.dumps(quotas), ex=7200)

emit("quiz\_started", {"quiz\_id": quiz\_id, "user\_id": user.sub, "tenant\_id":

payload.tenant\_id, "mode": mode, "filters": payload.filters.model\_dump(), "quota":  
quotas})

qids = list({v["q"] for v in vcache}[:payload.filters.num\_questions])

return QuizCreated(quiz\_id=quiz\_id, question\_ids=qids, expires\_at=expires\_at,  
mode=mode)

```

@router.get("/{quiz_id}/next", response_model=NextQuestion,
dependencies=[Depends(require_roles("student","admin"))])
def next_question(quiz_id: str, db: Session = Depends(get_db)):
    raw = redis_client.get(_rk(quiz_id,"versions"))
    if not raw: raise HTTPException(404, "Quiz not found or expired")
    versions = json.loads(raw)

    curk = _rk(quiz_id, "cursor")
    cur = int(redis_client.get(curk) or 0)
    if cur >= len(versions): raise HTTPException(404, "No more questions")

    quota = json.loads(redis_client.get(_rk(quiz_id),"quota_remaining") or "{}")

    # candidate window + hydrate IRT + SH p
    window = versions[cur : min(cur+20, len(versions))]
    candidates = []
    for w in window:
        ic =
        db.scalar(select(ItemCalibration).where(ItemCalibration.question_id==w["q"],
        ItemCalibration.version==w["v"])).limit(1)
        exp =
        db.scalar(select(ItemExposureControl).where(ItemExposureControl.question_id
        ==w["q"], ItemExposureControl.version==w["v"])).limit(1))
        a = (ic.a if ic and ic.a is not None else 1.0) if ic else 1.0
        b = (ic.b if ic and ic.b is not None else 0.0) if ic else 0.0
        c = (ic.c if ic and ic.c is not None else 0.2) if ic else 0.2
        sh_p = exp.sh_p if exp else 1.0
        candidates.append({"question_id": w["q"], "version": w["v"], "topic_id": w["t"],
        "a": a, "b": b, "c": c, "sh_p": sh_p})

    best = select_with_SH_and_quota(candidates, theta=0.0,
    remaining_quota=quota) or candidates[0]
    # decrement quota
    tkey = str(best.get("topic_id"))
    if tkey in quota and quota[tkey] > 0: quota[tkey] -= 1
    redis_client.set(_rk(quiz_id,"quota_remaining"), json.dumps(quota))
    redis_client.set(curk, cur+1)

    qv =
    db.scalar(select(QuestionVersion).where(QuestionVersion.question_id==best["que
    stion_id"], QuestionVersion.version==best["version"]))
    if not qv: raise HTTPException(500, "Item not found")
    opts =

```

```

db.execute(select(QuestionOption).where(QuestionOption.question_version_id==
qv.id)).scalars().all()
    bump_exposure(best["question_id"], best["version"])

```

```

    payload = {"stem_md": qv.stem_md, "lead_in": qv.lead_in, "options": [{"label":
o.option_label, "text": o.option_text_md} for o in opts]}
    emit("question_served", {"quiz_id": quiz_id, "question_id": best["question_id"],
"version": best["version"]})
    return NextQuestion(question_id=best["question_id"], version=best["version"],
payload=payload)

```

```

@router.post("/{quiz_id}/answers", response_model=AnswerResult,
dependencies=[Depends(require_roles("student","admin"))])
def submit_answer(quiz_id: str, payload: AnswerSubmit, db: Session =
Depends(get_db)):
    qv =
db.scalar(select(QuestionVersion).where(QuestionVersion.question_id==payload.q
uestion_id).order_by(QuestionVersion.version.desc()))
    if not qv: raise HTTPException(404, "Question not found")
    opts =
db.execute(select(QuestionOption).where(QuestionOption.question_version_id==
qv.id)).scalars().all()
    correct = next((o.option_label for o in opts if o.is_correct), None)
    if not correct: raise HTTPException(500, "No correct option set")
    ok = (payload.selected.upper() == correct)
    emit("answer_submitted", {"quiz_id": quiz_id, "question_id":
payload.question_id, "version": qv.version, "selected": payload.selected,
"is_correct": ok, "time_taken_ms": payload.time_taken_ms or 0})
    return AnswerResult(correct=ok, correct_option=correct,
explanation={"rationale_md": qv.rationale_md}, difficulty=0.5)
"""

```

```

# ----- Backend: API main + minor deps -----
write(f"{ROOT}/qbank-backend/app/main.py", """
from fastapi import FastAPI
from fastapi.middleware.cors import CORSMiddleware
from app.api.quizzes import router as quizzes_router
from app.api.author import router as author_router
from app.api.auth import router as auth_router

```

```

app = FastAPI(title="QBank API v4", version="4.0.0")
app.add_middleware(CORSMiddleware, allow_origins=["*"],
allow_credentials=True, allow_methods=["*"], allow_headers=["*"])

```

```

app.include_router(auth_router, prefix="/v1/auth", tags=["auth"])
app.include_router(quizzes_router, prefix="/v1/quizzes", tags=["quizzes"])
app.include_router(author_router, prefix="/v1/author", tags=["authoring"])

@app.get("/health")
def health(): return {"status": "ok"}
"""

# ----- Backend: core files (reusing v3 shapes) -----
write(f"{ROOT}/qbank-backend/app/core/config.py", """
import os
from dotenv import load_dotenv
load_dotenv()

DATABASE_URL = os.getenv("DATABASE_URL", "postgresql+psycopg2://
qbank:qbank@localhost:5432/qbank")
REDIS_URL = os.getenv("REDIS_URL", "redis://localhost:6379/0")
KAFKA_BOOTSTRAP = os.getenv("KAFKA_BOOTSTRAP", "localhost:9092")
KAFKA_TOPIC_EVENTS = os.getenv("KAFKA_TOPIC_EVENTS", "events.qbank")
TENANT_ID = os.getenv("APP_TENANT_ID",
"00000000-0000-0000-0000-000000000001")
APP_SECRET = os.getenv("APP_SECRET", "dev-secret-change-me")
MAX_DAILY_EXPOSURES = int(os.getenv("MAX_DAILY_EXPOSURES", "500"))
""")

write(f"{ROOT}/qbank-backend/app/core/database.py", """
from sqlalchemy import create_engine
from sqlalchemy.orm import sessionmaker
from app.core.config import DATABASE_URL

engine = create_engine(DATABASE_URL, future=True, pool_pre_ping=True)
SessionLocal = sessionmaker(bind=engine, autocommit=False, autoflush=False,
future=True)

def get_db():
    db = SessionLocal()
    try:
        yield db
    finally:
        db.close()
""")

write(f"{ROOT}/qbank-backend/app/core/cache.py", """

```

```

import redis
from datetime import datetime
from app.core.config import REDIS_URL, MAX_DAILY_EXPOSURES

redis_client = redis.Redis.from_url(REDIS_URL, decode_responses=True)

def exposure_key(question_id: int, version: int) -> str:
    day = datetime.utcnow().strftime("%Y%m%d")
    return f"exp:{day}:{question_id}:{version}"

def can_serve(question_id: int, version: int) -> bool:
    key = exposure_key(question_id, version)
    count = int(redis_client.get(key) or 0)
    return count < MAX_DAILY_EXPOSURES

def bump_exposure(question_id: int, version: int) -> None:
    key = exposure_key(question_id, version)
    pipe = redis_client.pipeline()
    pipe.incr(key, 1)
    pipe.expire(key, 86400) # 1 day
    pipe.execute()
    """)

write(f"{ROOT}/qbank-backend/app/core/events.py", """
import json
from uuid import uuid4
from datetime import datetime, timezone
from kafka import KafkaProducer
from app.core.config import KAFKA_BOOTSTRAP, KAFKA_TOPIC_EVENTS

producer = KafkaProducer(
    bootstrap_servers=KAFKA_BOOTSTRAP,
    value_serializer=lambda v: json.dumps(v).encode("utf-8"),
    key_serializer=lambda k: (k or "").encode("utf-8"),
    linger_ms=50, acks='1'
)

def emit(event_type: str, payload: dict):
    env = {"event_id": str(uuid4()), "event_type": event_type, "timestamp":
datetime.now(timezone.utc).isoformat(), **payload}
    producer.send(KAFKA_TOPIC_EVENTS, key=env["event_id"], value=env)
    """)

write(f"{ROOT}/qbank-backend/app/core/auth.py", """

```



```

from fastapi import Depends, HTTPException, status
from fastapi.security import HTTPBearer, HTTPAuthorizationCredentials
from pydantic import BaseModel
from typing import List
import jwt
from datetime import datetime, timedelta, timezone
from app.core.config import APP_SECRET

class TokenData(BaseModel):
    sub: str
    roles: List[str]

bearer = HTTPBearer()

def create_token(user_id: str, roles: List[str], ttl_minutes: int = 120) -> str:
    now = datetime.now(timezone.utc)
    payload = {"sub": user_id, "roles": roles, "iat": int(now.timestamp()), "exp":
int((now + timedelta(minutes=ttl_minutes)).timestamp())}
    return jwt.encode(payload, APP_SECRET, algorithm="HS256")

def get_current_user(creds: HTTPAuthorizationCredentials = Depends(bearer)) ->
TokenData:
    try:
        payload = jwt.decode(creds.credentials, APP_SECRET,
algorithms=["HS256"])
        return TokenData(sub=payload["sub"], roles=payload.get("roles", []))
    except Exception:
        raise HTTPException(status_code=status.HTTP_401_UNAUTHORIZED,
detail="Invalid or expired token")

def require_roles(*required: str):
    def checker(user: TokenData = Depends(get_current_user)):
        roles = set(user.roles)
        if not roles.intersection(set(required)):
            raise HTTPException(status_code=403, detail="Insufficient role")
        return user
    return checker

"""
write(f"{ROOT}/qbank-backend/app/api/auth.py", """
from fastapi import APIRouter
from pydantic import BaseModel
from typing import List
from app.core.auth import create_token

```

```

router = APIRouter()

class MockLogin(BaseModel):
    user_id: str
    roles: List[str]

@router.post("/mock-login")
def mock_login(payload: MockLogin):
    token = create_token(payload.user_id, payload.roles)
    return {"access_token": token, "token_type": "bearer", "roles": payload.roles}
"""

# ----- Authoring API (reuse v3 minimal) -----
write(f"{ROOT}/qbank-backend/app/api/author.py", """
from fastapi import APIRouter, Depends, HTTPException
from pydantic import BaseModel, constr
from typing import List, Optional
from sqlalchemy.orm import Session
from sqlalchemy import select, func

from app.core.database import get_db
from app.core.config import TENANT_ID
from app.core.auth import require_roles, TokenData
from app.models.orm import Topic, Question, QuestionVersion, QuestionOption,
QuestionPublication

router = APIRouter()

class OptionIn(BaseModel):
    label: constr(min_length=1, max_length=1)
    text_md: str
    is_correct: bool

class QuestionCreate(BaseModel):
    external_ref: Optional[str] = None
    topic_name: str
    exam_code: str = "DEMO-EXAM"
    stem_md: str
    lead_in: str
    rationale_md: str
    difficulty_label: Optional[str] = "medium"
    options: List[OptionIn]

```

```

@router.post("/questions",
dependencies=[Depends(require_roles("author","admin"))])
def create_question(payload: QuestionCreate, user: TokenData =
Depends(require_roles("author","admin")), db: Session = Depends(get_db)):
    t = db.scalar(select(Topic)×where(Topic×name == payload.topic_name))
    if not t:
        t = Topic(tenant_id=TENANT_ID, parent_id=None,
name=payload.topic_name, blueprint_code=None)
        db.add(t); db.flush()

    q = Question(tenant_id=TENANT_ID, external_ref=payload.external_ref,
created_by=user.sub, is_deleted=False)
    db.add(q); db.flush()

    next_v = (db.scalar(select(func.coalesce(func.max(QuestionVersion.version),
0)).where(QuestionVersion.question_id == q.id)) or 0) + 1
    qv = QuestionVersion(
        question_id=q.id, version=next_v, state="published",
        stem_md=payload.stem_md, lead_in=payload.lead_in,
rationale_md=payload.rationale_md,
        difficulty_label=payload×difficulty_label, topic_id=txid, tags={}, assets=[],
references=[]
    )
    db.add(qv); db.flush()

    for o in payload.options:
        db.add(QuestionOption(question_version_id=qv.id,
option_label=o.label.upper(), option_text_md=o.text_md,
is_correct=o.is_correct))

    db.add(QuestionPublication(question_id=q.id, live_version=next_v,
exam_code=payload.exam_code, tenant_id=TENANT_ID))
    db.commit()
    return {"question_id": q.id, "version": next_v, "topic_id": t.id}

@router.post("/publish/{question_id}",
dependencies=[Depends(require_roles("publisher","admin"))])
def publish(question_id: int, exam_code: str = "DEMO-EXAM", db: Session =
Depends(get_db)):
    qv = db.scalar(select(QuestionVersion)×where(QuestionVersion×question_id
== question_id).order_by(QuestionVersion.version.desc()))
    if not qv:
        raise HTTPException(404, "Question not found")

```

```

    pub = QuestionPublication(question_id=question_id, live_version=qv.version,
exam_code=exam_code, tenant_id=TENANT_ID)
    db.add(pub); db.commit()
    return {"published": True, "question_id": question_id, "version": qv.version,
"exam_code": exam_code}
""")

```

# ----- Student UI (Next.js) -----

```

write(f"{ROOT}/student-ui/package.json", ""

```

```

{
  "name": "student-ui",
  "private": true,
  "version": "0.1.0",
  "scripts": { "dev": "next dev", "build": "next build", "start": "next start" },
  "dependencies": { "next": "14.2.5", "react": "18.2.0", "react-dom": "18.2.0" }
}

```

```

""")

```

```

write(f"{ROOT}/student-ui/next.config.js", "module.exports = { reactStrictMode:
true };\\n")

```

```

write(f"{ROOT}/student-ui/.env.local", "NEXT_PUBLIC_API=http://
localhost:8000\\n")

```

```

write(f"{ROOT}/student-ui/pages/index.tsx", ""

```

```

import { useState } from 'react';

```

```

import Router from 'next/router';

```

```

export default function Home() {

```

```

  const [token, setToken] = useState('');

```

```

  const [tenant, setTenant] =

```

```

useState('00000000-0000-0000-0000-000000000001');

```

```

  const [exam, setExam] = useState('DEMO-EXAM');

```

```

  const start = async () => {

```

```

    sessionStorage.setItem('token', token);

```

```

    sessionStorage.setItem('tenant', tenant);

```

```

    sessionStorage.setItem('exam', exam);

```

```

    Router.push('/quiz');

```

```

  };

```

```

  return (

```

```

    <main style={{padding:24, maxWidth:640}}>

```

```

      <h1>QBank Student UI</h1>

```

```

      <p>Paste your JWT (from <code>/v1/auth/mock-login</code>)</p>

```

```

      <textarea value={token} onChange={(e)=>setToken(e.target.value)} rows={5}

```

```

style={{width:'100%'}} />
    <div style={{marginTop:12}}>
        <label>Tenant ID:&nbsp;</label>
        <input value={tenant} onChange={(e)=>setTenant(e.target.value)}
style={{width:'100%'}} />
    </div>
    <div style={{marginTop:12}}>
        <label>Exam Code:&nbsp;</label>
        <input value={exam} onChange={(e)=>setExam(e.target.value)}
style={{width:'100%'}} />
    </div>
    <button onClick={start} style={{marginTop:16}}>Start Quiz</button>
</main>
);
}
""")

```

```

write(f"{ROOT}/student-ui/pages/quiz.tsx", ""
import { useEffect, useState } from 'react';

```

```

const API = process.env.NEXT_PUBLIC_API || 'http://localhost:8000';

```

```

type NextQ = { question_id:number, version:number, payload:{ stem_md:string,
lead_in:string, options:{label:string,text:string}[] } };

```

```

export default function Quiz() {
    const [quizId, setQuizId] = useState<string>('');
    const [q, setQ] = useState<NextQ | null>(null);
    const [msg, setMsg] = useState<string>('');
    const token = typeof window !== 'undefined' ? sessionStorage.getItem('token') ||
'' : '';
    const tenant = typeof window !== 'undefined' ? sessionStorage.getItem('tenant')
|| '' : '';
    const exam = typeof window !== 'undefined' ? sessionStorage.getItem('exam') ||
'DEMO-EXAM' : 'DEMO-EXAM';

```

```

    const headers = { 'Content-Type':'application/json', 'Authorization': `Bearer $
{token}` };

```

```

    const createQuiz = async () => {
        const r = await fetch(`${API}/v1/quizzes`, {
            method: 'POST', headers,
            body: JSON.stringify({ tenant_id: tenant, adaptive: true, filters:
{ num_questions: 40, mode: 'tutor', exam_code: exam } })

```

```

});
const data = await r.json();
setQuizId(data.quiz_id);
};

const next = async () => {
  if (!quizId) return;
  const r = await fetch(`${API}/v1/quizzes/${quizId}/next`, { headers });
  if (r.ok) setQ(await r.json());
  else setMsg('No more questions.');
```

```

};

const answer = async (label: string) => {
  if (!quizId || !q) return;
  const r = await fetch(`${API}/v1/quizzes/${quizId}/answers`, {
    method: 'POST', headers,
    body: JSON.stringify({ question_id: q.question_id, selected: label,
time_taken_ms: 10000 })
  });
  const data = await r.json();
  setMsg(data.correct ? '✅ Correct' : '❌ Incorrect (Answer: $
{data.correct_option})');
  setTimeout(()=>{ setMsg(''); next(); }, 1000);
};

useEffect(()=>{ createQuiz().then(()=> setTimeout(next, 300)); }, []);

return (
  <main style={{padding:24, maxWidth:720}}>
    <h1>Quiz</h1>
    <p><b>Quiz ID:</b> {quizId}</p>
    {q ? (
      <div style={{border:'1px solid #ddd', borderRadius:8, padding:16}}>
        <div style={{whiteSpace:'pre-wrap'}}>{q.payload.stem_md}</div>
        <p><i>{q.payload.lead_in}</i></p>
        <div>
          {q.payload.options.map(o => (
            <button key={o.label} onClick={()=>answer(o.label)}
style={{display:'block', margin:'8px 0', width:'100%'}}>
              {o.label}. {o.text}
            </button>
          ))}
        </div>
      </div>
    )}
  </div>
)

```

```

    ) : <p>Loading...</p>{
    {msg && <p style={{marginTop:12}}>{msg}</p>}
  </main>
);
}
""")

```

```

# ----- Looker (LookML) over dbt marts -----
write(f"{ROOT}/looker/qbank.model.lkml", ""
connection: "qbank_postgres"

```

```

include: "views/*.view.lkml"

```

```

explore: fact_responses { }
explore: cohort_performance { }
""")

```

```

write(f"{ROOT}/looker/views/fact_responses.view.lkml", ""
view: fact_responses {
  sql_table_name: analytics.fact_responses ;;

```

```

  dimension: user_id { primary_key: yes; type: string; sql: ${TABLE}.user_id ;; }
  dimension: question_id { type: number; sql: ${TABLE}.question_id ;; }
  dimension: version { type: number; sql: ${TABLE}.version ;; }
  dimension: topic_id { type: number; sql: ${TABLE}.topic_id ;; }
  dimension_group: ts {
    type: time
    timeframes: [date, week, month, year]
    sql: ${TABLE}.ts ;;
  }
  measure: attempts { type: count; drill_fields: [user_id, question_id, version] }
  measure: accuracy { type: average; sql: CASE WHEN ${TABLE}.is_correct THEN 1
ELSE 0 END ;; value_format_name: "percent_2" }
  measure: avg_time_ms { type: average; sql: ${TABLE}.time_taken_ms ;; }
}
""")

```

```

write(f"{ROOT}/looker/views/cohort_performance.view.lkml", ""
view: cohort_performance {
  sql_table_name: analytics.cohort_performance ;;

```

```

  dimension_group: day { type: time; timeframes: [date, week, month, year]; sql: $
${TABLE}.day ;; }
  dimension: cohort { type: string; sql: ${TABLE}.cohort ;; }

```

```

    measure: attempts { type: sum; sql: ${TABLE}.attempts ;; }
    measure: accuracy { type: average; sql: ${TABLE}.accuracy ;;
value_format_name: "percent_2" }
}
""")

```

```

# ----- Metabase setup (API import script + sample cards) -----
write(f"{ROOT}/metabase/README.md", "")
# Metabase Setup (against Postgres analytics schema)

```

- 1) Connect your Metabase instance to the same Postgres DB that dbt populates (`analytics` schema).
- 2) Optionally import the sample dashboard and cards via API:

```

```bash
export MB_URL="http://localhost:3000"
export MB_USER="admin@example.com"
export MB_PASS="yourpass"
export MB_DB_ID=2 # The database ID Metabase assigned to your Postgres
source
bash import.sh
""")

write(f"{ROOT}/metabase/import.sh", "")
#!/usr/bin/env bash
set -euo pipefail
if [ -z "${MB_URL:-}" ] || [ -z "${MB_USER:-}" ] || [ -z "${MB_PASS:-}" ] || [ -z "${MB_DB_ID:-}" ]; then
echo "Set MB_URL, MB_USER, MB_PASS, MB_DB_ID"; exit 1; fi

```

```

TOKEN=$(curl -s -X POST "$MB_URL/api/session" -H "Content-Type: application/
json" -d '{"username":"$MB_USER","password":"$MB_PASS"}' | jq -r .id)

```

```

create_card () {
local payload="$1"
curl -s -X POST "$MB_URL/api/card" -H "Content-Type: application/json" -H "X-
Metabase-Session: $TOKEN" -d "$payload" >/dev/null
}

```

```

FACT_RESP_CARD='{
"name": "Accuracy by Day",
"dataset_query": {
"type": "native",
"native": { "query": "select ts::date as day, avg(case when is_correct then 1 else 0

```



```
end) as accuracy from analytics.fact_responses group by 1 order by 1 desc" },
"database": "$MB_DB_ID"
},
"display": "line"
}'
```

```
COHORT_CARD='{
"name": "Cohort Attempts by Day",
"dataset_query": {
"type": "native",
"native": { "query": "select day, cohort, attempts from
analytics.cohort_performance order by day desc" },
"database": "$MB_DB_ID"
},
"display": "bar"
}'
```

```
create_card "$FACT_RESP_CARD"
create_card "$COHORT_CARD"
echo "Imported sample cards."
""")
```

----- CI: GitHub Actions -----

```
write(f"{ROOT}/.github/workflows/backend.yml", ""
name: backend-ci
on:
push: { paths: ["qbank-backend/", ".github/workflows/backend.yml"] }
pull_request: { paths: ["qbank-backend/", ".github/workflows/backend.yml"] }
jobs:
api:
runs-on: ubuntu-latest
services:
postgres:
image: postgres:16-alpine
env:
POSTGRES_USER: qbank
POSTGRES_PASSWORD: qbank
POSTGRES_DB: qbank
ports: ["5432:5432"]
options: >-
--health-cmd="pg_isready -U qbank" --health-interval=10s --health-timeout=5s
--health-retries=5
redis:
image: redis:7-alpine
```

```

ports: ["6379:6379"]
env:
  DATABASE_URL: postgresql+psycpg2://qbank:qbank@localhost:5432/qbank
  REDIS_URL: redis://localhost:6379/0
  APP_SECRET: ci-secret
steps:
  - uses: actions/checkout@v4
  - name: Set up Python
    uses: actions/setup-python@v5
    with: { python-version: "3.11" }
  - name: Install deps
    run: |
      pip install -r qbank-backend/requirements.txt
  - name: Lint
    run: |
      python -m black --check qbank-backend
      flake8 qbank-backend
  - name: Alembic upgrade head
    working-directory: qbank-backend
    run: |
      alembic -c alembic.ini upgrade head
    """)

```

```

write(f"{ROOT}/.github/workflows/student-ui.yml", """
name: student-ui
on:
  push: { paths: ["student-ui/", ".github/workflows/student-ui.yml"] }
  pull_request: { paths: ["student-ui/", ".github/workflows/student-ui.yml"] }
jobs:
  build:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v4
      - uses: actions/setup-node@v4
        with: { node-version: "20" }
      - name: Install & build
        working-directory: student-ui
        run: |
          npm ci || npm install
          npm run build
    """)

```

```

----- Zip -----
zip_path = "/mnt/data/qbank_scaffold_v4.zip"

```

```
with zipfile.ZipFile(zip_path, "w", zipfile.ZIP_DEFLATED) as z:
for p in pathlib.Path(ROOT).rglob("*"):
z.write(p, p.relative_to(ROOT))
```

zip\_path

Always show details

Copy

Result

'/mnt/data/qbank\_scaffold\_v4.zip'

Shipped! I added everything you asked for and wrapped it in a fresh bundle:

Download the scaffold (v4)

What's new in v4 (quick scan):

Alembic autogenerate wired to Base.metadata + a sample migration file; ready for alembic revision --autogenerate.

CI (GitHub Actions):

backend-ci: lints, spins up Postgres/Redis services, runs alembic upgrade head.

student-ui: installs & builds the Next.js app.

Sympson-Hetter exposure control:

New table item\_exposure\_control (question\_id, version, sh\_p).

Selector samples  $u \sim U(0,1)$  and serves only if  $u \leq sh\_p$ , otherwise goes to the next best item.

Student UI (Next.js):

Paste JWT → create quiz → see question → answer → next.

Analytics models:

Looker: looker/qbank.model.lkml + views for fact\_responses and cohort\_performance.

Metabase: metabase/import.sh to POST sample cards (accuracy over time; cohort attempts).

How to try the new bits fast:

Run DB migration two ways

SQL: apply /sql/item\_exposure\_control.sql along with your existing DDLs

or Alembic: cd qbank-backend && alembic -c alembic.ini upgrade head

Turn on SH for a few items

scss

Copy

Edit

```
INSERT INTO item_exposure_control(question_id, version, sh_p) VALUES  
(123,1,0.3),(456,2,0.5);
```

Student UI

cd student-ui && npm install && npm run dev → http://localhost:3000

Get a token via POST /v1/auth/mock-login and paste it on the UI homepage.

CI

Commit .github/workflows/\* and push – the backend job will bring up Postgres and run migrations.